

```
In [292... import torch
from torch import nn
import torch.optim as optim
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
import time
from sklearn.model_selection import KFold
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import torch.nn.functional as F
```

## Problem 1

```
In [53]: train_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\fashion-mnist_train.csv"
test_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\fashion-mnist_test.csv"
train_numpy = np.loadtxt(train_path, dtype = np.float32, delimiter = ",", skiprows=1)
test_numpy = np.loadtxt(test_path, dtype = np.float32, delimiter = ",", skiprows=1)
print(train_numpy.shape)
```

(60000, 785)

```
In [75]: def try_gpu(i=0):
    if torch.cuda.device_count() >= i+1:
        return torch.device(f'cuda:{i}')
    return torch.device('cpu')
```

```
In [226... def train_model(model, total_train, total_test, loss_fn, optim, epochs, update_freq):
    train_loss_hist = []
    train_acc_hist = []
    test_acc_hist = []

    main_tic = time.perf_counter()
    for epoch in range(1, epochs + 1):
        tic = time.perf_counter()
        loss_train = 0
        correct_train = 0
        for imgs in train_loader:
            labels = imgs[:, -1].long()
            imgs = imgs[:, :-1]

            batch_size = imgs.shape[0]
            outputs = model(imgs.view(batch_size, -1))
            _, predicted = torch.max(outputs, dim = 1)
            correct_train += int((predicted == labels).sum())
            loss = loss_fn(outputs, labels)

            optim.zero_grad()
            loss.backward()
            optim.step()

            loss_train += loss.item()

        toc = time.perf_counter()
```

```

correct_val = 0
with torch.no_grad():
    for imgs in val_loader:
        labels = imgs[:, -1].long()
        imgs = imgs[:, :-1]
        batch_size = imgs.shape[0]
        outputs = model(imgs.view(batch_size, -1))
        _, predicted = torch.max(outputs, dim = 1)
        correct_val += int((predicted == labels).sum())

    if epoch == 1 or epoch == epochs or epoch % update_freq == 0:
        print(f"Epoch {epoch}:\n\tDuration = {round(toc - tic, 3)} seconds\n\tTrain Loss = {round(loss_train, 3)}\n\tTrain Accuracy = {round(correct_train/total_train, 3)}\n\tTest Accuracy = {round(correct_val/total_test, 3)}")

    train_loss_hist.append(round(loss_train / len(train_loader), 5))
    train_acc_hist.append(round(correct_train/total_train, 3))
    test_acc_hist.append(round(correct_val/total_test, 3))

main_toc = time.perf_counter()
print(f"\nTotal Training Time = {round(main_toc - main_tic, 3)} seconds\nAverage Training Time per Epoch = {round(main_toc - main_tic, 3) / epochs} seconds")
return train_loss_hist, train_acc_hist, test_acc_hist

```

```

In [216... def plot_model(title, fig_num, loss_hist, train_hist, test_hist, leg_loc):
    plot = plt.figure(fig_num)
    x = range(1, len(loss_hist)+1)
    plt.plot(x, loss_hist)
    plt.plot(x, train_hist)
    plt.plot(x, test_hist)
    plt.legend(["Training Loss", "Training Accuracy", "Testing Accuracy"], loc=leg_loc)
    plt.xlabel('Epoch')
    plt.ylabel('Error/Accuracy')
    plt.title(title)

```

```

In [228... train_tensor = torch.from_numpy(train_numpy)
test_tensor = torch.from_numpy(test_numpy)

train_features = train_tensor[:, 1:]
test_features = test_tensor[:, 1:]

train_label = torch.unsqueeze(train_tensor[:, 0], dim=-1)
test_label = torch.unsqueeze(test_tensor[:, 0], dim=-1)

train_set = torch.cat((train_features, train_label), dim=1).to(device=try_gpu())
test_set = torch.cat((test_features, test_label), dim=1).to(device=try_gpu())

train_loader = torch.utils.data.DataLoader(train_set, batch_size=600, shuffle=True)
val_loader = torch.utils.data.DataLoader(test_set, batch_size=test_set.shape[0], shuffle=False)

```

```

In [229... model_1 = nn.Sequential(
    nn.Identity(), # Using Identity Layers as place-holders for Dropout Layers
    nn.Linear(784, 1024), # First Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(1024, 512), # Second Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(512, 256), # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(256, 10)).to(device=try_gpu()) # Output Layer

```

```
optimizer_1 = optim.SGD(model_1.parameters(), lr=1e-3)

model_1.eval()
```

```
Out[229]: Sequential(
  (0): Identity()
  (1): Linear(in_features=784, out_features=1024, bias=True)
  (2): ReLU()
  (3): Identity()
  (4): Linear(in_features=1024, out_features=512, bias=True)
  (5): ReLU()
  (6): Identity()
  (7): Linear(in_features=512, out_features=256, bias=True)
  (8): ReLU()
  (9): Linear(in_features=256, out_features=10, bias=True)
)
```

```
In [ ]: # class Module_1a(nn.Module):
#       def __init__(self):
#           super().__init__()
#           self.fc1 = nn.Linear(784, 1024)
#           self.act1 = nn.ReLU()
#           self.fc2 = nn.Linear(1024, 512)
#           self.act2 = nn.ReLU()
#           self.fc3 = nn.Linear(512, 256)
#           self.act3 = nn.ReLU()
#           self.fc4 = nn.Linear(256, 10)

#       def forward(self, x):
#           out = self.act1(self.fc1(x))
#           out = self.act2(self.fc2(out))
#           out = self.act3(self.fc3(out))
#           out = self.fc4(out)
#           return out
```

```
In [230...] t_loss_hist_a, t_acc_hist_a, v_acc_hist_a = train_model(model_1,
                                                                    train_set.shape[0],
                                                                    test_set.shape[0],
                                                                    nn.CrossEntropyLoss(),
                                                                    optimizer_1,
                                                                    20,
                                                                    5)
```

Epoch 1:  
Duration = 1.662 seconds  
Training Loss: 1.08243  
Training Accuracy: 0.719  
Validation Accuracy: 0.8

Epoch 5:  
Duration = 1.523 seconds  
Training Loss: 0.43935  
Training Accuracy: 0.847  
Validation Accuracy: 0.847

Epoch 10:  
Duration = 1.544 seconds  
Training Loss: 0.37876  
Training Accuracy: 0.865  
Validation Accuracy: 0.864


Epoch 15:  
Duration = 1.58 seconds  
Training Loss: 0.34601  
Training Accuracy: 0.878  
Validation Accuracy: 0.869

Epoch 20:  
Duration = 1.563 seconds  
Training Loss: 0.321  
Training Accuracy: 0.886  
Validation Accuracy: 0.871

Total Training Time = 34.137 seconds  
Average Training Time per Epoch = 1.707 seconds

In [231...

```
title_1 = "Figure 1 - Loss and Accuracy per Epoch for Model 1a"  
plot_model(title_1, 1, t_loss_hist_a, t_acc_hist_a, v_acc_hist_a, 'upper right')
```



```
In [232...] torch.save(model_1.state_dict(), "homework_0-model_1.params")
```

```
In [234...] model_2 = nn.Sequential(
    nn.Identity(),
    nn.Linear(784, 1024),    # First Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(1024, 512),    # Second Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(512, 256),    # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(256, 10)).to(device=try_gpu()) # Output Layer

optimizer_2 = optim.SGD(model_2.parameters(), lr = 1e-3, weight_decay=1e-4)

model_2.eval()
```

```
Out[234]: Sequential(
  (0): Identity()
  (1): Linear(in_features=784, out_features=1024, bias=True)
  (2): ReLU()
  (3): Identity()
  (4): Linear(in_features=1024, out_features=512, bias=True)
  (5): ReLU()
  (6): Identity()
  (7): Linear(in_features=512, out_features=256, bias=True)
  (8): ReLU()
  (9): Linear(in_features=256, out_features=10, bias=True)
)
```

```
In [235...] t_loss_hist_b, t_acc_hist_b, v_acc_hist_b = train_model(model_2,
                                                                    train_set.shape[0],
                                                                    test_set.shape[0],
                                                                    nn.CrossEntropyLoss(),
                                                                    optimizer_2,
                                                                    20,
                                                                    5)
```

Epoch 1:  
Duration = 1.694 seconds  
Training Loss: 1.06484  
Training Accuracy: 0.714  
Validation Accuracy: 0.795

Epoch 5:  
Duration = 1.852 seconds  
Training Loss: 0.44831  
Training Accuracy: 0.843  
Validation Accuracy: 0.844

Epoch 10:  
Duration = 1.772 seconds  
Training Loss: 0.3828  
Training Accuracy: 0.865  
Validation Accuracy: 0.855

Epoch 15:  
Duration = 1.693 seconds  
Training Loss: 0.34954  
Training Accuracy: 0.876  
Validation Accuracy: 0.864

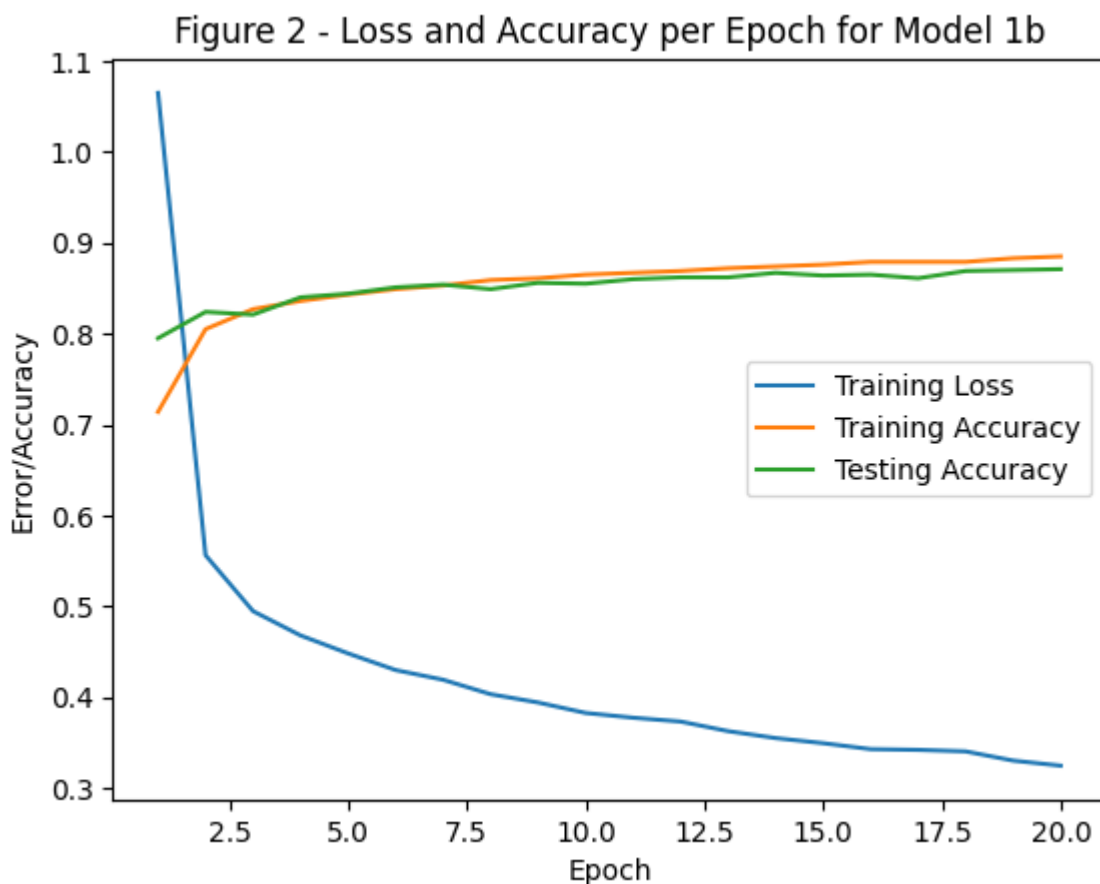
Epoch 20:  
Duration = 1.829 seconds  
Training Loss: 0.32481  
Training Accuracy: 0.885  
Validation Accuracy: 0.871

Total Training Time = 37.19 seconds

Average Training Time per Epoch = 1.86 seconds

In [236...

```
title_2 = "Figure 2 - Loss and Accuracy per Epoch for Model 1b"  
plot_model(title_2, 2, t_loss_hist_b, t_acc_hist_b, v_acc_hist_b, 'right')
```



```
In [237...] model_3 = nn.Sequential(
    nn.Dropout(p=0.3),
    nn.Linear(784, 1024),    # First Hidden Layer
    nn.ReLU(),
    nn.Dropout(p=0.3),
    nn.Linear(1024, 512),   # Second Hidden Layer
    nn.ReLU(),
    nn.Dropout(p=0.3),
    nn.Linear(512, 256),    # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(256, 10)).to(device=try_gpu()) # Output Layer

optimizer_3 = optim.SGD(model_3.parameters(), lr = 1e-3)

model_3.eval()
```

```
Out[237]: Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): Linear(in_features=784, out_features=1024, bias=True)
  (2): ReLU()
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=1024, out_features=512, bias=True)
  (5): ReLU()
  (6): Dropout(p=0.3, inplace=False)
  (7): Linear(in_features=512, out_features=256, bias=True)
  (8): ReLU()
  (9): Linear(in_features=256, out_features=10, bias=True)
)
```

```
In [238...] t_loss_hist_c, t_acc_hist_c, v_acc_hist_c = train_model(model_3,
                                                                    train_set.shape[0],
                                                                    test_set.shape[0],
                                                                    nn.CrossEntropyLoss(),
                                                                    optimizer_3,
                                                                    20,
                                                                    5)
```

Epoch 1:  
Duration = 2.05 seconds  
Training Loss: 1.23049  
Training Accuracy: 0.71  
Validation Accuracy: 0.784

Epoch 5:  
Duration = 1.537 seconds  
Training Loss: 0.44936  
Training Accuracy: 0.841  
Validation Accuracy: 0.846

Epoch 10:  
Duration = 1.564 seconds  
Training Loss: 0.38318  
Training Accuracy: 0.864  
Validation Accuracy: 0.863

Epoch 15:  
Duration = 1.516 seconds  
Training Loss: 0.34995  
Training Accuracy: 0.876  
Validation Accuracy: 0.867


Epoch 20:  
Duration = 1.525 seconds  
Training Loss: 0.32561  
Training Accuracy: 0.884  
Validation Accuracy: 0.872

Total Training Time = 34.5 seconds

Average Training Time per Epoch = 1.725 seconds

In [246...

```
title_3 = "Figure 3 - Loss and Accuracy per Epoch for Model 1c"  
plot_model(title_3, 3, t_loss_hist_c, t_acc_hist_c, v_acc_hist_c, 'upper right')
```





```
In [240... model_4 = model_2  
model_4.load_state_dict(torch.load("homework_0-model_1.params"))
```

Out[240]: <All keys matched successfully>

```
In [241... t_loss_hist_d, t_acc_hist_d, v_acc_hist_d = train_model(model_4,  
                                                         train_set.shape[0],  
                                                         test_set.shape[0],  
                                                         nn.CrossEntropyLoss(),  
                                                         optimizer_2,  
                                                         20,  
                                                         5)
```

Epoch 1:

Duration = 1.678 seconds  
Training Loss: 0.31526  
Training Accuracy: 0.888  
Validation Accuracy: 0.871

Epoch 5:

Duration = 1.532 seconds  
Training Loss: 0.30065  
Training Accuracy: 0.894  
Validation Accuracy: 0.877

Epoch 10:

Duration = 1.607 seconds  
Training Loss: 0.28236  
Training Accuracy: 0.9  
Validation Accuracy: 0.875

Epoch 15:

Duration = 1.687 seconds  
Training Loss: 0.26986  
Training Accuracy: 0.903  
Validation Accuracy: 0.88

Epoch 20:


Duration = 1.579 seconds  
Training Loss: 0.25747  
Training Accuracy: 0.91  
Validation Accuracy: 0.883

Total Training Time = 35.706 seconds

Average Training Time per Epoch = 1.785 seconds

```
In [242... title_4 = "Figure 4 - Loss and Accuracy per Epoch for Model 1b with Preloaded Weights"  
plot_model(title_4, 4, t_loss_hist_d, t_acc_hist_d, v_acc_hist_d, 'right')
```

Figure 4 - Loss and Accuracy per Epoch for Model 1b with Preloaded Weights



```
In [243...] model_5 = model_3
model_5.load_state_dict(torch.load("homework_0-model_1.params"))
```

```
Out[243]: <All keys matched successfully>
```

```
In [244...] t_loss_hist_e, t_acc_hist_e, v_acc_hist_e = train_model(model_5,
                                                                train_set.shape[0],
                                                                test_set.shape[0],
                                                                nn.CrossEntropyLoss(),
                                                                optimizer_3,
                                                                20,
                                                                5)
```

Epoch 1:  
Duration = 1.554 seconds  
Training Loss: 0.31538  
Training Accuracy: 0.888  
Validation Accuracy: 0.875

Epoch 5:  
Duration = 1.441 seconds  
Training Loss: 0.29839  
Training Accuracy: 0.894  
Validation Accuracy: 0.873

Epoch 10:  
Duration = 1.567 seconds  
Training Loss: 0.28498  
Training Accuracy: 0.9  
Validation Accuracy: 0.882

Epoch 15:  
Duration = 1.555 seconds  
Training Loss: 0.2721  
Training Accuracy: 0.904  
Validation Accuracy: 0.876

Epoch 20:  
Duration = 1.506 seconds  
Training Loss: 0.25842  
Training Accuracy: 0.909  
Validation Accuracy: 0.876


Total Training Time = 34.356 seconds

Average Training Time per Epoch = 1.718 seconds

In [245...

```
title_5 = "Figure 5 - Loss and Accuracy per Epoch for Model 1c with Preloaded Weights"  
plot_model(title_5, 5, t_loss_hist_e, t_acc_hist_e, v_acc_hist_e, 'right')
```

Figure 5 - Loss and Accuracy per Epoch for Model 1c with Preloaded Weights



## Problem 2

```
In [271... house_train_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\housing_train.csv"
house_test_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\housing_test.csv"
house_raw_train = pd.read_csv(house_train_path, delimiter = ",", header=0)
house_raw_test = pd.read_csv(house_test_path, delimiter = ",", header=0)
```

```
In [272... label = 'SalePrice'
features = pd.concat([house_raw_train.drop(columns=['Id',label]),
                      house_raw_test.drop(columns=['Id'])])

numeric_features = features.dtypes[features.dtypes != 'object'].index
features[numeric_features] = features[numeric_features].apply(
    lambda x: (x - x.mean()) / (x.std()))

features[numeric_features] = features[numeric_features].fillna(0)
features = pd.get_dummies(features, dummy_na=True)

house_train_dataframe = features[:house_raw_train.shape[0]].copy()
house_train_dataframe[label] = house_raw_train[label]
house_test_dataframe = features[house_raw_train.shape[0]:].copy()
```

```
In [544... def loss_1(predictions, labels):
    outputs = torch.square(torch.log10(labels) - torch.log10(torch.squeeze(predictions)))
    output = torch.sqrt(torch.mean(outputs))
    # print(predictions.shape, labels.shape, predictions, labels)
    # print(torch.squeeze(predictions).shape, outputs, output)
    return output.requires_grad_()
```

```
In [550... def plot_model_kfold(title, fig_num, train_loss, test_loss, leg_loc):
    plot = plt.figure(fig_num)
    x = range(1, len(train_loss)+1)
    plt.plot(x, train_loss)
    plt.plot(x, test_loss)
    plt.legend(["Training Loss", "Testing Loss"], loc=leg_loc)
    plt.xlabel('Epoch')
    plt.ylabel('Error')
    plt.title(title)
```

```
In [311... house_train_tensor = torch.tensor(house_train_dataframe.values).to(device=try_gpu()).t
house_test_tensor = torch.tensor(house_test_dataframe.values).to(device=try_gpu()).to
```

```
In [545... def train_model_kfold(model, train_tensor, batch, loss_fn, optim, epochs, update_freq,
                        k_folds):
    kf = KFold(n_splits=k_folds, shuffle=True, random_state=0)
    main_tic = time.perf_counter()
    train_loss_hist = []
    test_loss_hist = []

    for epoch in range(1, epochs + 1):
        train_loss_epoch = []
        test_loss_epoch = []
        tic = time.perf_counter()

        for i, (train_index, test_index) in enumerate(kf.split(train_tensor)):
            train, test = train_tensor[train_index], train_tensor[test_index]
```

```

t_loader = torch.utils.data.DataLoader(train, batch_size=batch, shuffle=True)
v_loader = torch.utils.data.DataLoader(test, batch_size=test.shape[0], shuffle=False)
loss_train = 0
loss_val = 0

for imgs in t_loader:
    labels = imgs[:, -1]
    imgs = imgs[:, :-1]
    batch_size = imgs.shape[0]
    outputs = model(imgs.view(batch_size, -1))
    loss_t = loss_fn(outputs, labels)
    optim.zero_grad()
    loss_t.backward()
    optim.step()
    loss_train += loss_t.item()

with torch.no_grad():
    for imgs in v_loader:
        labels = imgs[:, -1]
        imgs = imgs[:, :-1]
        batch_size = imgs.shape[0]
        outputs = model(imgs.view(batch_size, -1))
        loss_v = loss_fn(outputs, labels)
        loss_val += loss_v.item()

train_loss_epoch.append(round(4*loss_train/len(train), 5))
test_loss_epoch.append(round(loss_val, 5))

toc = time.perf_counter()
train_loss_hist.append(sum(train_loss_epoch)/len(train_loss_epoch))
test_loss_hist.append(sum(test_loss_epoch)/len(test_loss_epoch))

if epoch == 1 or epoch == epochs or epoch % update_freq == 0:
    print(f"Epoch {epoch}:\n\tDuration = {round(toc - tic, 3)} seconds\n\tAverage Train Loss = {round(train_loss_hist[-1], 3)}\n\tAverage Test Loss = {round(test_loss_hist[-1], 3)}")

main_toc = time.perf_counter()
print(f"\nTotal Training Time = {round(main_toc - main_tic, 3)} seconds\nAverage Test Loss = {round(test_loss_hist[-1], 3)}")
return train_loss_hist, test_loss_hist

```

In [548...

```

model_house_1 = nn.Sequential(
    nn.Identity(),
    nn.LazyLinear(512),    # First Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(512, 256),  # Second Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(256, 128),  # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(128, 1)).to(device=try_gpu()) # Output Layer

optimizer_house_1 = optim.SGD(model_house_1.parameters(), lr = 1e-3)

model_house_1.eval()

```

```
Out[548]: Sequential(
  (0): Identity()
  (1): LazyLinear(in_features=0, out_features=512, bias=True)
  (2): ReLU()
  (3): Identity()
  (4): Linear(in_features=512, out_features=256, bias=True)
  (5): ReLU()
  (6): Identity()
  (7): Linear(in_features=256, out_features=128, bias=True)
  (8): ReLU()
  (9): Linear(in_features=128, out_features=1, bias=True)
)
```

```
In [549... t_loss_hist_house_a, v_loss_hist_house_a = train_model_kfold(model_house_1,
                                                                house_train_tensor,
                                                                4,
                                                                loss_1,
                                                                optimizer_house_1,
                                                                20,
                                                                1,
                                                                5)
```

Epoch 1:  
Duration = 2.776 seconds  
Average Training Loss: 5.02613  
Average Validation Loss: 4.89212

Epoch 2:  
Duration = 2.525 seconds  
Average Training Loss: 4.14203  
Average Validation Loss: 4.08028

Epoch 3:  
Duration = 2.583 seconds  
Average Training Loss: 3.64616  
Average Validation Loss: 3.60849

Epoch 4:  
Duration = 2.533 seconds  
Average Training Loss: 3.32055  
Average Validation Loss: 3.29435

Epoch 5:  
Duration = 2.514 seconds  
Average Training Loss: 3.08113  
Average Validation Loss: 3.06126

Epoch 6:  
Duration = 2.52 seconds  
Average Training Loss: 2.89232  
Average Validation Loss: 2.87653

Epoch 7:  
Duration = 2.831 seconds  
Average Training Loss: 2.7368  
Average Validation Loss: 2.72374

Epoch 8:  
Duration = 2.695 seconds  
Average Training Loss: 2.6046  
Average Validation Loss: 2.59356

Epoch 9:  
Duration = 2.652 seconds  
Average Training Loss: 2.48965  
Average Validation Loss: 2.48022

Epoch 10:  
Duration = 2.669 seconds  
Average Training Loss: 2.38798  
Average Validation Loss: 2.37988

Epoch 11:  
Duration = 2.724 seconds  
Average Training Loss: 2.29704  
Average Validation Loss: 2.28989

Epoch 12:  
Duration = 3.121 seconds  
Average Training Loss: 2.21465  
Average Validation Loss: 2.20833

Epoch 13:  
Duration = 3.249 seconds  
Average Training Loss: 2.13945  
Average Validation Loss: 2.13376

Epoch 14:  
Duration = 2.779 seconds  
Average Training Loss: 2.07004  
Average Validation Loss: 2.06508

Epoch 15:  
Duration = 2.823 seconds  
Average Training Loss: 2.00594  
Average Validation Loss: 2.00144

Epoch 16:  
Duration = 3.841 seconds  
Average Training Loss: 1.94625  
Average Validation Loss: 1.94216

Epoch 17:  
Duration = 3.622 seconds  
Average Training Loss: 1.89023  
Average Validation Loss: 1.88668

Epoch 18:  
Duration = 3.915 seconds  
Average Training Loss: 1.83785  
Average Validation Loss: 1.83454


Epoch 19:  
Duration = 3.348 seconds  
Average Training Loss: 1.78829  
Average Validation Loss: 1.78537

Epoch 20:  
Duration = 3.083 seconds  
Average Training Loss: 1.74152  
Average Validation Loss: 1.73886

Total Training Time = 58.804 seconds  
Average Training Time per Epoch = 2.94 seconds

In [559...

```
title_6 = "Figure 6 - Regression Loss for Model 2b"  
plot_model_kfold(title_6, 6, t_loss_hist_house_a, v_loss_hist_house_a, 'right')
```



In [555...

```
model_house_2 = nn.Sequential(  
    nn.Dropout(p=0.3),  
    nn.LazyLinear(512), # First Hidden Layer  
    nn.ReLU(),
```



```

nn.Dropout(p=0.3),
nn.Linear(512, 256),    # Second Hidden Layer
nn.ReLU(),
nn.Dropout(p=0.3),
nn.Linear(256, 128),    # Third Hidden Layer
nn.ReLU(),
nn.Linear(128, 1)).to(device=try_gpu()) # Output Layer

optimizer_house_2 = optim.SGD(model_house_2.parameters(), lr = 1e-3, weight_decay=1e-4)

model_house_2.eval()

```

Out[555]:

```

Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): LazyLinear(in_features=0, out_features=512, bias=True)
  (2): ReLU()
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=512, out_features=256, bias=True)
  (5): ReLU()
  (6): Dropout(p=0.3, inplace=False)
  (7): Linear(in_features=256, out_features=128, bias=True)
  (8): ReLU()
  (9): Linear(in_features=128, out_features=1, bias=True)
)

```

In [556...]

```

t_loss_hist_house_b, v_loss_hist_house_b = train_model_kfold(model_house_2,
                                                                house_train_tensor,
                                                                4,
                                                                loss_1,
                                                                optimizer_house_2,
                                                                20,
                                                                1,
                                                                5)

```

Epoch 1:  
Duration = 2.86 seconds  
Average Training Loss: 4.94739  
Average Validation Loss: 4.80709

Epoch 2:  
Duration = 2.836 seconds  
Average Training Loss: 4.07012  
Average Validation Loss: 4.01156

Epoch 3:  
Duration = 2.81 seconds  
Average Training Loss: 3.59933  
Average Validation Loss: 3.56339

Epoch 4:  
Duration = 3.096 seconds  
Average Training Loss: 3.28708  
Average Validation Loss: 3.26182

Epoch 5:  
Duration = 2.879 seconds  
Average Training Loss: 3.05532  
Average Validation Loss: 3.03613

Epoch 6:  
Duration = 3.168 seconds  
Average Training Loss: 2.87162  
Average Validation Loss: 2.85622

Epoch 7:  
Duration = 3.148 seconds  
Average Training Loss: 2.71951  
Average Validation Loss: 2.70683

Epoch 8:  
Duration = 2.804 seconds  
Average Training Loss: 2.58987  
Average Validation Loss: 2.57918

Epoch 9:  
Duration = 2.976 seconds  
Average Training Loss: 2.47702  
Average Validation Loss: 2.4678

Epoch 10:  
Duration = 2.938 seconds  
Average Training Loss: 2.37706  
Average Validation Loss: 2.36903

Epoch 11:  
Duration = 2.851 seconds  
Average Training Loss: 2.28738  
Average Validation Loss: 2.28034

Epoch 12:  
Duration = 2.778 seconds  
Average Training Loss: 2.20612  
Average Validation Loss: 2.19987

Epoch 13:  
Duration = 2.918 seconds  
Average Training Loss: 2.13176  
Average Validation Loss: 2.12623

Epoch 14:  
Duration = 2.819 seconds  
Average Training Loss: 2.0633  
Average Validation Loss: 2.05837

Epoch 15:  
Duration = 3.014 seconds  
Average Training Loss: 1.9998  
Average Validation Loss: 1.99545

Epoch 16:  
Duration = 2.861 seconds  
Average Training Loss: 1.94076  
Average Validation Loss: 1.9368

Epoch 17:  
Duration = 2.787 seconds  
Average Training Loss: 1.88549  
Average Validation Loss: 1.8819

Epoch 18:  
Duration = 2.991 seconds  
Average Training Loss: 1.8335  
Average Validation Loss: 1.83029

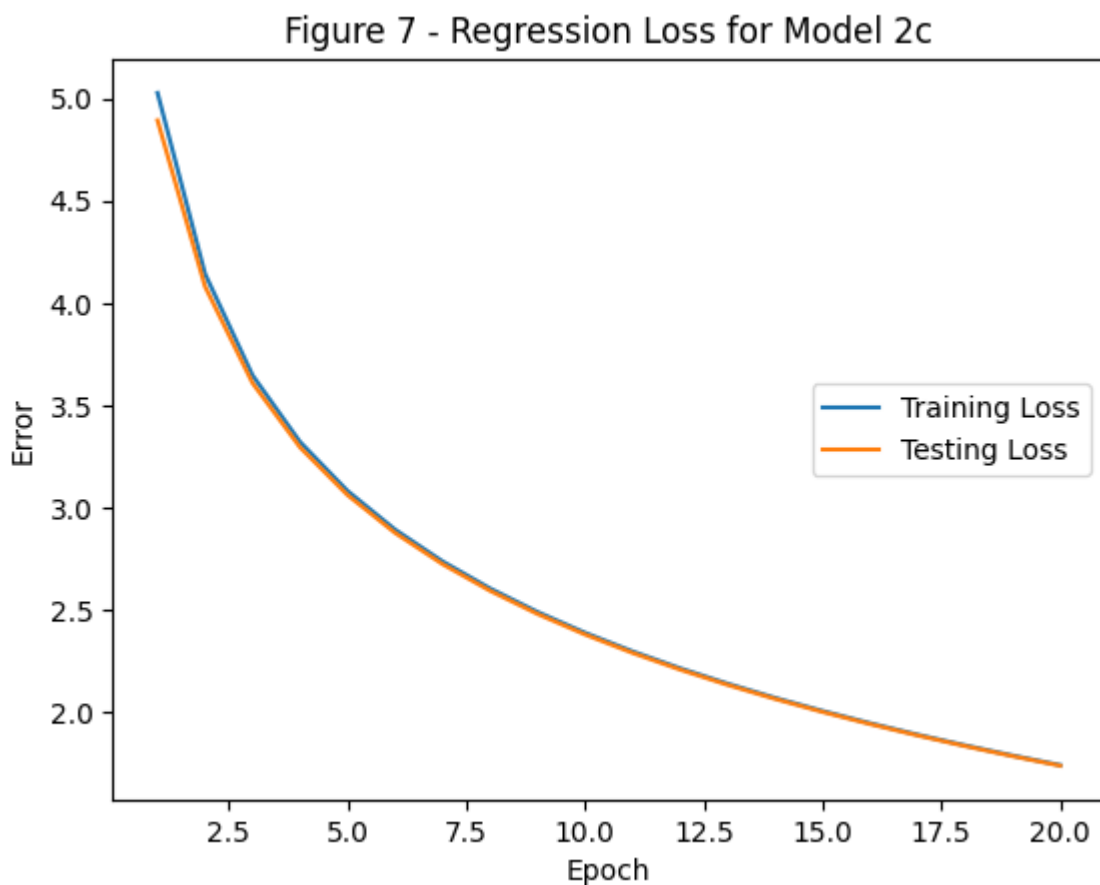
Epoch 19:  
Duration = 3.275 seconds  
Average Training Loss: 1.78446  
Average Validation Loss: 1.78159

Epoch 20:  
Duration = 3.338 seconds  
Average Training Loss: 1.73803  
Average Validation Loss: 1.73552

Total Training Time = 59.147 seconds  
Average Training Time per Epoch = 2.957 seconds

In [560...

```
title_7 = "Figure 7 - Regression Loss for Model 2c"  
plot_model_kfold(title_7, 7, t_loss_hist_house_a, v_loss_hist_house_a, 'right')
```



In [580...

```
preds = model_house_2(house_test_tensor)  
ensemble_preds = torch.exp(torch.squeeze(preds))  
submission = pd.DataFrame({'Id':house_raw_test.Id,
```

```
        'SalePrice':ensemble_preds.detach().cpu().numpy()})  
submission.to_csv('submission.csv', index=False)
```

```
torch.Size([1459, 1])  
torch.Size([1459])
```