

[https://github.com/Christian-Martens-UNCC/ECGR-4106/tree/main/Homework\\_0-Multi\\_Layer\\_Perceptron](https://github.com/Christian-Martens-UNCC/ECGR-4106/tree/main/Homework_0-Multi_Layer_Perceptron)

Problem 1:

- A) It does not appear that my network needs more epochs for full training as the training and validation accuracy have converged and are pretty high. There does appear to be some very slight overfitting at the very end of the training since the training accuracy is climbing slightly while the testing accuracy is pretty much stationary.

```
Epoch 1:
  Duration = 1.662 seconds
  Training Loss: 1.08243
  Training Accuracy: 0.719
  Validation Accuracy: 0.8


Epoch 5:
  Duration = 1.523 seconds
  Training Loss: 0.43935
  Training Accuracy: 0.847
  Validation Accuracy: 0.847

Epoch 10:
  Duration = 1.544 seconds
  Training Loss: 0.37876
  Training Accuracy: 0.865
  Validation Accuracy: 0.864

Epoch 15:
  Duration = 1.58 seconds
  Training Loss: 0.34601
  Training Accuracy: 0.878
  Validation Accuracy: 0.869

Epoch 20:
  Duration = 1.563 seconds
  Training Loss: 0.321
  Training Accuracy: 0.886
  Validation Accuracy: 0.871

Total Training Time = 34.137 seconds
Average Training Time per Epoch = 1.707 seconds
```



- B) The training results are very similar to the baseline. The validation and training accuracies are more inconsistent but there also seems to be less overfitting compared to the baseline model. During other training sessions, this method also resulted in a ~2% increase in validation accuracy.

```
In [234]: model_2 = nn.Sequential(
    nn.Identity(),
    nn.Linear(784, 1024), # First Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(1024, 512), # Second Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(512, 256), # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(256, 10)).to(device=try_gpu()) # Output Layer

optimizer_2 = optim.SGD(model_2.parameters(), lr = 1e-3, weight_decay=1e-4)

model_2.eval()
```

Out[234]: Sequential(

- (0): Identity()
- (1): Linear(in\_features=784, out\_features=1024, bias=True)
- (2): ReLU()
- (3): Identity()
- (4): Linear(in\_features=1024, out\_features=512, bias=True)
- (5): ReLU()
- (6): Identity()
- (7): Linear(in\_features=512, out\_features=256, bias=True)
- (8): ReLU()
- (9): Linear(in\_features=256, out\_features=10, bias=True)

Epoch 1:  
Duration = 1.694 seconds  
Training Loss: 1.06484  
Training Accuracy: 0.714  
Validation Accuracy: 0.795


Epoch 5:  
Duration = 1.852 seconds  
Training Loss: 0.44831  
Training Accuracy: 0.843  
Validation Accuracy: 0.844

Epoch 10:  
Duration = 1.772 seconds  
Training Loss: 0.3828  
Training Accuracy: 0.865  
Validation Accuracy: 0.855

Epoch 15:  
Duration = 1.693 seconds  
Training Loss: 0.34954  
Training Accuracy: 0.876  
Validation Accuracy: 0.864

Epoch 20:  
Duration = 1.829 seconds  
Training Loss: 0.32481  
Training Accuracy: 0.885  
Validation Accuracy: 0.871

Total Training Time = 37.19 seconds  
Average Training Time per Epoch = 1.86 seconds



- C) The baseline model and the weight decay model both performed similarly to the dropout model. This method was consistently faster than the other two however, likely due to the reduced number of computations that needed completing. During other training sessions, this method generally performed with the highest validation accuracy, sometimes reaching as high as 90%.

```

In [237]: model_3 = nn.Sequential(
            nn.Dropout(p=0.3),
            nn.Linear(784, 1024), # First Hidden Layer
            nn.ReLU(),
            nn.Dropout(p=0.3),
            nn.Linear(1024, 512), # Second Hidden Layer
            nn.ReLU(),
            nn.Dropout(p=0.3),
            nn.Linear(512, 256), # Third Hidden Layer
            nn.ReLU(),
            nn.Linear(256, 10)).to(device=try_gpu()) # Output Layer

optimizer_3 = optim.SGD(model_3.parameters(), lr = 1e-3)
model_3.eval()

Out[237]: Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): Linear(in_features=784, out_features=1024, bias=True)
  (2): ReLU()
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=1024, out_features=512, bias=True)
  (5): ReLU()
  (6): Dropout(p=0.3, inplace=False)
  (7): Linear(in_features=512, out_features=256, bias=True)
  (8): ReLU()
  (9): Linear(in_features=256, out_features=10, bias=True)
)

```

Epoch 1:  
Duration = 2.05 seconds  
Training Loss: 1.23049  
Training Accuracy: 0.71  
Validation Accuracy: 0.784


Epoch 5:  
Duration = 1.537 seconds  
Training Loss: 0.44936  
Training Accuracy: 0.841  
Validation Accuracy: 0.846

Epoch 10:  
Duration = 1.564 seconds  
Training Loss: 0.38318  
Training Accuracy: 0.864  
Validation Accuracy: 0.863

Epoch 15:  
Duration = 1.516 seconds  
Training Loss: 0.34995  
Training Accuracy: 0.876  
Validation Accuracy: 0.867



Epoch 20:  
Duration = 1.525 seconds  
Training Loss: 0.32561  
Training Accuracy: 0.884  
Validation Accuracy: 0.872

Total Training Time = 34.5 seconds  
Average Training Time per Epoch = 1.725 seconds




- D) The training times for these both models were not much better than the untrained models. As we can see, there was not much change between the weights of the baseline model and the weights of the weight decay and dropout models. This method did allow us to further train our model, which was used to confirm that training had completed and that the models was slowly drifting towards overfitting if they continued for additional epochs.

Figure 4 - Loss and Accuracy per Epoch for Model 1b with Preloaded Weights Figure 5 - Loss and Accuracy per Epoch for Model 1c with Preloaded Weights



Problem 2:

- A) When we standardize the continuous numerical features of a dataset, the features are condensed such that all the features are relatively the same size compared to one another. This makes sure that the model doesn't get dominated by features with large input values and allows features with smaller input values to contribute to training the model. This overall helps with generalization.
- B) The model complexity was increased by adding in more hidden layers with more neurons. It is hard to compare training times since the model in class was ran on a different computer, but it appears that my model has a lower validation loss than the model in the lecture. I would expect the model to take longer to train given it has more layers and thus more computations.



```
In [548]: model_house_1 = nn.Sequential(
    nn.Identity(),
    nn.LazyLinear(512), # First Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(512, 256), # Second Hidden Layer
    nn.ReLU(),
    nn.Identity(),
    nn.Linear(256, 128), # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(128, 1)).to(device='try_gpu()') # Output Layer

optimizer_house_1 = optim.SGD(model_house_1.parameters(), lr = 1e-3)

model_house_1.eval()

Out[548]: Sequential(
  (0): Identity()
  (1): LazyLinear(in_features=0, out_features=512, bias=True)
  (2): ReLU()
  (3): Identity()
  (4): Linear(in_features=512, out_features=256, bias=True)
  (5): ReLU()
  (6): Identity()
  (7): Linear(in_features=256, out_features=128, bias=True)
  (8): ReLU()
  (9): Linear(in_features=128, out_features=1, bias=True)
)
```

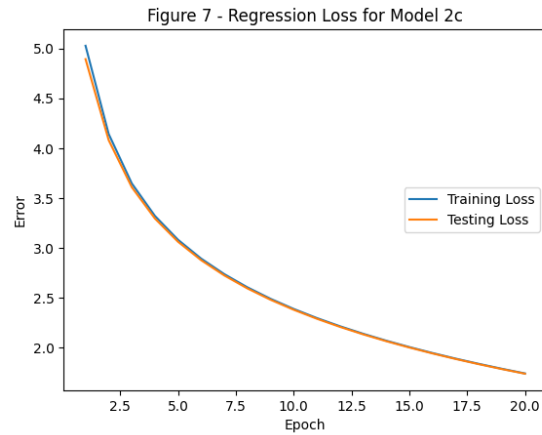
- C) This model trained significantly faster than the previous model (roughly a 10% decrease in training time). Additionally, we can see that these methods did not increase the validation loss.

```
In [555]: model_house_2 = nn.Sequential(
    nn.Dropout(p=0.3),
    nn.LazyLinear(512), # First Hidden Layer
    nn.ReLU(),
    nn.Dropout(p=0.3),
    nn.Linear(512, 256), # Second Hidden Layer
    nn.ReLU(),
    nn.Dropout(p=0.3),
    nn.Linear(256, 128), # Third Hidden Layer
    nn.ReLU(),
    nn.Linear(128, 1)).to(device='try_gpu()') # Output Layer

optimizer_house_2 = optim.SGD(model_house_2.parameters(), lr = 1e-3, weight_decay=1e-4)

model_house_2.eval()

Out[555]: Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): LazyLinear(in_features=0, out_features=512, bias=True)
  (2): ReLU()
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=512, out_features=256, bias=True)
  (5): ReLU()
  (6): Dropout(p=0.3, inplace=False)
  (7): Linear(in_features=256, out_features=128, bias=True)
  (8): ReLU()
  (9): Linear(in_features=128, out_features=1, bias=True)
)
```



D) My predictions to Kaggle were not as good as I would've liked. Had I more time, I would've gone back and tried other ways of implementing Pytorch so that the regression is more accurate. My predictions were all very high compared to what I was expecting.