

```
In [257... import torch
from torch import nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor
import time
from sklearn.model_selection import KFold
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import torch.nn.functional as F
import gc
from ptflops import get_model_complexity_info
```

Problem 1

```
In [86]: train_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\fashion-mnist_train.csv"
test_path = r"C:\Users\ccm51\OneDrive\Desktop\ECGR 4106\fashion-mnist_test.csv"
train_numpy = np.loadtxt(train_path, dtype = np.float32, delimiter = ",", skiprows=1)
test_numpy = np.loadtxt(test_path, dtype = np.float32, delimiter = ",", skiprows=1)
print(train_numpy.shape)
```

```
(60000, 785)
```

```
In [101... data_path = '../data-unversioned/ecgr4106/'
fashion_MNIST = datasets.FashionMNIST(data_path, train=True, download=True, transform=
fashion_MNIST_val = datasets.FashionMNIST(data_path, train=False, download=True, trans
```

```
In [123... len(fashion_MNIST_val)
```

```
Out[123]: 10000
```

```
In [5]: def try_gpu(i=0):
        if torch.cuda.device_count() >= i+1:
            return torch.device(f'cuda:{i}')
        return torch.device('cpu')
```

```
In [238... def training_loop(n_epochs, optimizer, model, loss_fn, train_loader, val_loader, update_fn):
    train_loss_hist = []
    train_acc_hist = []
    val_acc_hist = []
    main_tic = time.perf_counter()

    for epoch in range(1, n_epochs + 1):
        tic = time.perf_counter()
        loss_train = 0.0
        correct_train = 0
        correct_val = 0

        for imgs, labels in train_loader:
            # if torch.cuda.is_available():
            imgs = imgs.to(device=try_gpu())
            labels = labels.to(device=try_gpu())
            outputs = model(imgs)
```

```

        loss = loss_fn(outputs, labels)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        loss_train += loss.item()

    toc = time.perf_counter()

    with torch.no_grad():
        total = 0
        for imgs, labels in train_loader:
            outputs = model(imgs.to(device=try_gpu()))
            _, predicted = torch.max(outputs, dim=1)
            total += labels.shape[0]
            correct_train += int((predicted == labels.to(device=try_gpu())).sum())
        train_acc = round(correct_train/total, 3)
        total = 0
        for imgs, labels in val_loader:
            outputs = model(imgs.to(device=try_gpu()))
            _, predicted = torch.max(outputs, dim=1)
            total += labels.shape[0]
            correct_val += int((predicted == labels.to(device=try_gpu())).sum())
        val_acc = round(correct_val/total, 3)

    train_loss_hist.append(round(loss_train / len(train_loader), 5))
    train_acc_hist.append(train_acc)
    val_acc_hist.append(val_acc)

    if epoch == 1 or epoch == n_epochs or epoch % update_freq == 0:
        print(f"Epoch {epoch}:\n\tDuration = {round(toc - tic, 3)} seconds\n\tTrain Loss = {round(loss_train, 5)}\n\tTrain Accuracy = {train_acc}\n\tValidation Accuracy = {val_acc}")

    main_toc = time.perf_counter()
    print(f"\nTotal Training Time = {round(main_toc - main_tic, 3)} seconds\nAverage 1 Epoch Time = {round((main_toc - main_tic) / n_epochs, 3)} seconds")
    return train_loss_hist, train_acc_hist, val_acc_hist

```

```

In [7]: def plot_model(title, fig_num, loss_hist, train_hist, test_hist, leg_loc):
    plot = plt.figure(fig_num)
    x = range(1, len(loss_hist)+1)
    plt.plot(x, loss_hist)
    plt.plot(x, train_hist)
    plt.plot(x, test_hist)
    plt.legend(["Training Loss", "Training Accuracy", "Testing Accuracy"], loc=leg_loc)
    plt.xlabel('Epoch')
    plt.ylabel('Error/Accuracy')
    plt.title(title)

```

```

In [158... class alt1_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
        self.fc3 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

```

```

def forward(self, x):
    out = self.maxp1(self.relu(self.conv1(x)))
    out = self.maxp2(self.relu(self.conv2(out)))
    out = self.flat(out)
    out = self.relu(self.fc1(out))
    out = self.relu(self.fc2(out))
    out = self.fc3(out)
    return out

model_1 = alt1_LeNet().to(device=try_gpu())
optimizer_1 = optim.SGD(model_1.parameters(), lr=1e-1)
model_1.eval()

```

```

Out[158]: alt1_LeNet(
  (conv1): LazyConv2d(0, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 16, kernel_size=(5, 5), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [150... train_loader_1 = DataLoader(fashion_MNIST, batch_size=128, shuffle=True)
val_loader_1 = DataLoader(fashion_MNIST_val, batch_size=len(fashion_MNIST_val), shuffle=True)

```

```

Out[150]: 469

```

```

In [160... t_loss_hist_1, t_acc_hist_1, v_acc_hist_1 = training_loop(10,
                                                                optimizer_1,
                                                                model_1,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                1)

```


```
Epoch 1:
  Duration = 8.849 seconds
  Training Loss: 1.11996
  Training Accuracy: 0.739
  Validation Accuracy: 0.727
Epoch 2:
  Duration = 8.972 seconds
  Training Loss: 0.53999
  Training Accuracy: 0.83
  Validation Accuracy: 0.822
Epoch 3:
  Duration = 9.038 seconds
  Training Loss: 0.44165
  Training Accuracy: 0.859
  Validation Accuracy: 0.849
Epoch 4:
  Duration = 8.816 seconds
  Training Loss: 0.38818
  Training Accuracy: 0.858
  Validation Accuracy: 0.848
Epoch 5:
  Duration = 8.709 seconds
  Training Loss: 0.35384
  Training Accuracy: 0.867
  Validation Accuracy: 0.855
Epoch 6:
  Duration = 8.659 seconds
  Training Loss: 0.33247
  Training Accuracy: 0.869
  Validation Accuracy: 0.855
Epoch 7:
  Duration = 9.653 seconds
  Training Loss: 0.31426
  Training Accuracy: 0.889
  Validation Accuracy: 0.872
Epoch 8:
  Duration = 8.907 seconds
  Training Loss: 0.29852
  Training Accuracy: 0.892
  Validation Accuracy: 0.877
Epoch 9:
  Duration = 8.869 seconds
  Training Loss: 0.28339
  Training Accuracy: 0.898
  Validation Accuracy: 0.877
Epoch 10:
  Duration = 9.13 seconds
  Training Loss: 0.27235
  Training Accuracy: 0.899
  Validation Accuracy: 0.879
```

```
Total Training Time = 164.335 seconds
Average Training Time per Epoch = 16.433 seconds
```

In [162...

```
title_1 = "Figure 1 - Loss and Accuracy per Epoch for the updated LeNet Model"
plot_model(title_1, 1, t_loss_hist_1, t_acc_hist_1, v_acc_hist_1, 'upper right')
```

Figure 1 - Loss and Accuracy per Epoch for the updated LeNet Model



Problem 2

In [164...

```

class alt2a_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=3, padding=1)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=3)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
        self.fc3 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.fc3(out)
        return out

model_2a = alt2a_LeNet().to(device=try_gpu())
optimizer_2a = optim.SGD(model_2a.parameters(), lr=1e-1)
model_2a.eval()

```

```

Out[164]: alt2a_LeNet(
    (conv1): LazyConv2d(0, 6, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (conv2): LazyConv2d(0, 16, kernel_size=(3, 3), stride=(1, 1))
    (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
    (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
    (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
    (relu): ReLU()
    (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [166... t_loss_hist_2a, t_acc_hist_2a, v_acc_hist_2a = training_loop(10,
                                                                optimizer_2a,
                                                                model_2a,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                5)

```

Epoch 1:

```

    Duration = 8.57 seconds
    Training Loss: 1.0938
    Training Accuracy: 0.718
    Validation Accuracy: 0.708

```

Epoch 5:

```

    Duration = 9.154 seconds
    Training Loss: 0.35228
    Training Accuracy: 0.878
    Validation Accuracy: 0.868

```

Epoch 10:

```

    Duration = 8.199 seconds
    Training Loss: 0.27944
    Training Accuracy: 0.901
    Validation Accuracy: 0.886

```

```

Total Training Time = 157.307 seconds
Average Training Time per Epoch = 15.731 seconds


```

```

In [167... title_2a = "Figure 2a - Loss and Accuracy per Epoch for the altered LeNet Model (Shrir
plot_model(title_2a, 2, t_loss_hist_2a, t_acc_hist_2a, v_acc_hist_2a, 'upper right')

```

Figure 2a - Loss and Accuracy per Epoch for the altered LeNet Model (shrinking convolution window)



In [168...

```

class alt2b_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=7, padding=3)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=7)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
        self.fc3 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.fc3(out)
        return out

model_2b = alt2b_LeNet().to(device=try_gpu())
optimizer_2b = optim.SGD(model_2b.parameters(), lr=1e-1)
model_2b.eval()

```

```

Out[168]: alt2b_LeNet(
  (conv1): LazyConv2d(0, 6, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 16, kernel_size=(7, 7), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```
In [169... t_loss_hist_2b, t_acc_hist_2b, v_acc_hist_2b = training_loop(10,
                                                                optimizer_2b,
                                                                model_2b,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                5)
```

Epoch 1:

Duration = 9.395 seconds
 Training Loss: 1.08711
 Training Accuracy: 0.739
 Validation Accuracy: 0.73

Epoch 5:

Duration = 9.349 seconds
 Training Loss: 0.34582
 Training Accuracy: 0.875
 Validation Accuracy: 0.86

Epoch 10:


Duration = 10.162 seconds
 Training Loss: 0.27412
 Training Accuracy: 0.896
 Validation Accuracy: 0.876

Total Training Time = 186.064 seconds

Average Training Time per Epoch = 18.606 seconds

```
In [170... title_2b = "Figure 2b - Loss and Accuracy per Epoch for the altered LeNet Model (Expanding Convolution Window)"
plot_model(title_2b, 3, t_loss_hist_2b, t_acc_hist_2b, v_acc_hist_2b, 'upper right')
```

Figure 2b - Loss and Accuracy per Epoch for the altered LeNet Model (Expanding Convolution Window)



```
In [181... class alt3_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(12, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(32, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
```



```

self.fc3 = nn.LazyLinear(10)
self.relu = nn.ReLU()
self.flat = nn.Flatten()

def forward(self, x):
    out = self.maxp1(self.relu(self.conv1(x)))
    out = self.maxp2(self.relu(self.conv2(out)))
    out = self.flat(out)
    out = self.relu(self.fc1(out))
    out = self.relu(self.fc2(out))
    out = self.fc3(out)
    return out

model_3 = alt3_LeNet().to(device=try_gpu())
optimizer_3 = optim.SGD(model_3.parameters(), lr=1e-1)
model_3.eval()

```

```

Out[181]: alt3_LeNet(
  (conv1): LazyConv2d(0, 12, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 32, kernel_size=(5, 5), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [182... t_loss_hist_3, t_acc_hist_3, v_acc_hist_3 = training_loop(10,
                                                             optimizer_3,
                                                             model_3,
                                                             nn.CrossEntropyLoss(),
                                                             train_loader_1,
                                                             val_loader_1,
                                                             5)

```

```

Epoch 1:
  Duration = 11.009 seconds
  Training Loss: 1.02504
  Training Accuracy: 0.789
  Validation Accuracy: 0.78
Epoch 5:
  Duration = 10.568 seconds
  Training Loss: 0.32999
  Training Accuracy: 0.886
  Validation Accuracy: 0.871
Epoch 10:
  Duration = 11.317 seconds
  Training Loss: 0.24853
  Training Accuracy: 0.915
  Validation Accuracy: 0.896

```

```

Total Training Time = 190.937 seconds
Average Training Time per Epoch = 19.094 seconds

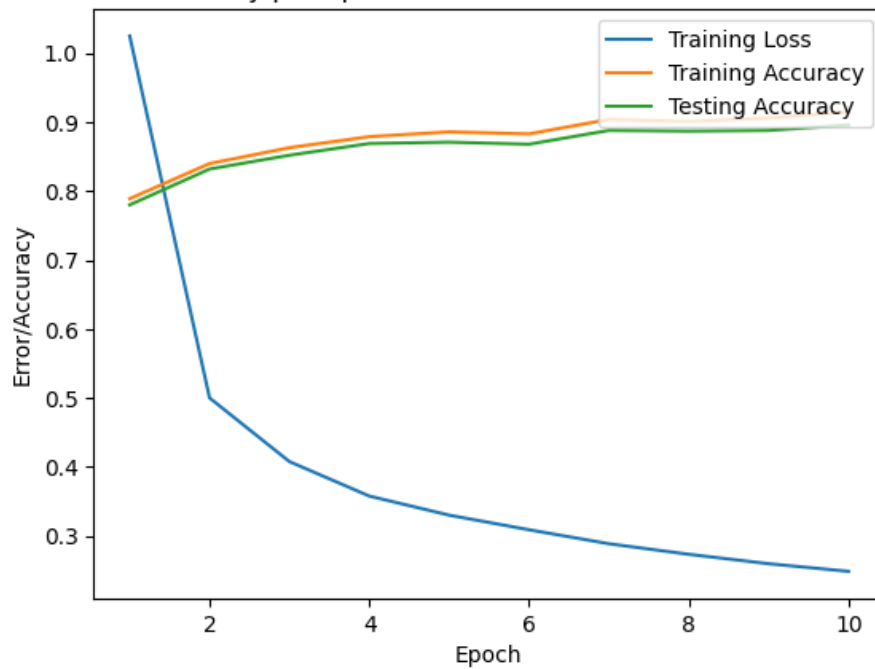
```

```

In [183... title_3 = "Figure 3 - Loss and Accuracy per Epoch for the altered LeNet Model (Doublir
plot_model(title_3, 4, t_loss_hist_3, t_acc_hist_3, v_acc_hist_3, 'upper right')

```

Figure 3 - Loss and Accuracy per Epoch for the altered LeNet Model (Doubling Layer Width)



In [196...

```

class alt4_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv3 = nn.LazyConv2d(8, kernel_size=5, padding=2)
        self.maxp3 = nn.MaxPool2d(kernel_size=2, stride=1)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
        self.fc3 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.maxp3(self.relu(self.conv3(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.fc3(out)
        return out

model_4 = alt4_LeNet().to(device=try_gpu())
optimizer_4 = optim.SGD(model_4.parameters(), lr=1e-1)
model_4.eval()

```

```

Out[196]: alt4_LeNet(
    (conv1): LazyConv2d(0, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (conv2): LazyConv2d(0, 16, kernel_size=(5, 5), stride=(1, 1))
    (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (conv3): LazyConv2d(0, 8, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (maxp3): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1, ceil_mode=False)
    (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
    (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
    (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
    (relu): ReLU()
    (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [197... t_loss_hist_4, t_acc_hist_4, v_acc_hist_4 = training_loop(10,
                                                                optimizer_4,
                                                                model_4,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                5)

```

Epoch 1:

```

    Duration = 9.411 seconds
    Training Loss: 1.42292
    Training Accuracy: 0.636
    Validation Accuracy: 0.627

```

Epoch 5:

```

    Duration = 9.522 seconds
    Training Loss: 0.37345
    Training Accuracy: 0.869
    Validation Accuracy: 0.856

```

Epoch 10:

```

    Duration = 9.648 seconds
    Training Loss: 0.29061
    Training Accuracy: 0.896
    Validation Accuracy: 0.876

```

Total Training Time = 172.741 seconds


Average Training Time per Epoch = 17.274 seconds

```

In [198... title_4 = "Figure 4 - Loss and Accuracy per Epoch for the altered LeNet Model (Adding
plot_model(title_4, 5, t_loss_hist_4, t_acc_hist_4, v_acc_hist_4, 'upper right')

```

Figure 4 - Loss and Accuracy per Epoch for the altered LeNet Model (Adding Convolution Layer)



In [199...

```

class alt5_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(256)
        self.fc2 = nn.Linear(120)
        self.fc3 = nn.Linear(84)
        self.fc4 = nn.Linear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.relu(self.fc3(out))
        out = self.fc4(out)
        return out

model_5 = alt5_LeNet().to(device=try_gpu())
optimizer_5 = optim.SGD(model_5.parameters(), lr=1e-1)
model_5.eval()

```

```

Out[199]: alt5_LeNet(
  (conv1): LazyConv2d(0, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 16, kernel_size=(5, 5), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=256, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc4): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [200... t_loss_hist_5, t_acc_hist_5, v_acc_hist_5 = training_loop(10,
                                                             optimizer_5,
                                                             model_5,
                                                             nn.CrossEntropyLoss(),
                                                             train_loader_1,
                                                             val_loader_1,
                                                             5)

```

Epoch 1:

```

    Duration = 9.023 seconds
    Training Loss: 1.3352
    Training Accuracy: 0.717
    Validation Accuracy: 0.709

```

Epoch 5:

```

    Duration = 9.593 seconds
    Training Loss: 0.36617
    Training Accuracy: 0.875
    Validation Accuracy: 0.863

```

Epoch 10:

```

    Duration = 11.919 seconds
    Training Loss: 0.27715
    Training Accuracy: 0.899
    Validation Accuracy: 0.882

```

Total Training Time = 195.642 seconds


Average Training Time per Epoch = 19.564 seconds

```

In [201... title_5 = "Figure 5 - Loss and Accuracy per Epoch for the altered LeNet Model (Adding
plot_model(title_5, 6, t_loss_hist_5, t_acc_hist_5, v_acc_hist_5, 'upper right')

```

Figure 5 - Loss and Accuracy per Epoch for the altered LeNet Model (Adding Fully Connected Layer)



In [204...

```

class alt6a_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(120)
        self.fc2 = nn.LazyLinear(84)
        self.fc3 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.fc3(out)
        return out

model_6a = alt6a_LeNet().to(device=try_gpu())
optimizer_6a = optim.SGD(model_6a.parameters(), lr=2e-1)
model_6a.eval()

```

```

Out[204]: alt6a_LeNet(
  (conv1): LazyConv2d(0, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 16, kernel_size=(5, 5), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```
In [205... t_loss_hist_6a, t_acc_hist_6a, v_acc_hist_6a = training_loop(10,
                                                                optimizer_6a,
                                                                model_6a,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                5)
```

Epoch 1:

Duration = 8.659 seconds
 Training Loss: 0.93004
 Training Accuracy: 0.826
 Validation Accuracy: 0.813

Epoch 5:

Duration = 11.621 seconds
 Training Loss: 0.30714
 Training Accuracy: 0.893
 Validation Accuracy: 0.878


Epoch 10:

Duration = 10.245 seconds
 Training Loss: 0.23812
 Training Accuracy: 0.918
 Validation Accuracy: 0.891

Total Training Time = 190.071 seconds
 Average Training Time per Epoch = 19.007 seconds

```
In [206... title_6a = "Figure 6a - Loss and Accuracy per Epoch for the altered LeNet Model (Doubl
plot_model(title_6a, 7, t_loss_hist_6a, t_acc_hist_6a, v_acc_hist_6a, 'upper right')
```

Figure 6a - Loss and Accuracy per Epoch for the altered LeNet Model (Doubling Learning Rate)



```
In [207... class alt6b_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(6, kernel_size=5, padding=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(16, kernel_size=5)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
```

```

self.fc1 = nn.LazyLinear(120)
self.fc2 = nn.LazyLinear(84)
self.fc3 = nn.LazyLinear(10)
self.relu = nn.ReLU()
self.flat = nn.Flatten()

def forward(self, x):
    out = self.maxp1(self.relu(self.conv1(x)))
    out = self.maxp2(self.relu(self.conv2(out)))
    out = self.flat(out)
    out = self.relu(self.fc1(out))
    out = self.relu(self.fc2(out))
    out = self.fc3(out)
    return out

model_6b = alt6b_LeNet().to(device=try_gpu())
optimizer_6b = optim.SGD(model_6b.parameters(), lr=5e-2)
model_6b.eval()

```

```

Out[207]: alt6b_LeNet(
  (conv1): LazyConv2d(0, 6, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 16, kernel_size=(5, 5), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [208...] t_loss_hist_6b, t_acc_hist_6b, v_acc_hist_6b = training_loop(10,
                                                                optimizer_6b,
                                                                model_6b,
                                                                nn.CrossEntropyLoss(),
                                                                train_loader_1,
                                                                val_loader_1,
                                                                5)

```

```

Epoch 1:
  Duration = 10.28 seconds
  Training Loss: 1.54733
  Training Accuracy: 0.702
  Validation Accuracy: 0.698

```

```

Epoch 5:
  Duration = 8.634 seconds
  Training Loss: 0.41036
  Training Accuracy: 0.858
  Validation Accuracy: 0.847

```

```

Epoch 10:
  Duration = 11.12 seconds
  Training Loss: 0.32063
  Training Accuracy: 0.889
  Validation Accuracy: 0.874

```

```

Total Training Time = 193.704 seconds
Average Training Time per Epoch = 19.37 seconds

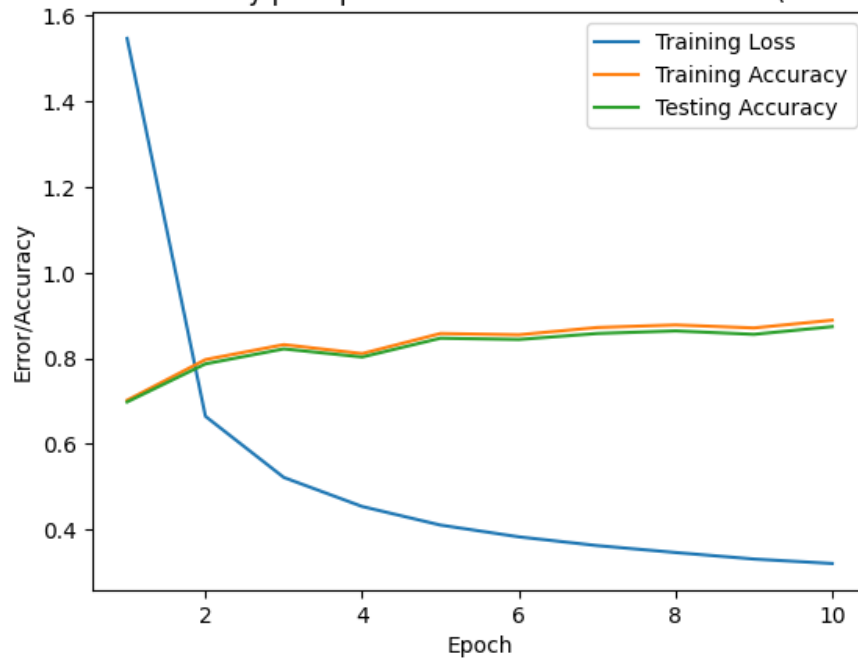
```

```

In [209...] title_6b = "Figure 6b - Loss and Accuracy per Epoch for the altered LeNet Model (Halvi
plot_model(title_6b, 8, t_loss_hist_6b, t_acc_hist_6b, v_acc_hist_6b, 'upper right')

```


Figure 6b - Loss and Accuracy per Epoch for the altered LeNet Model (Halving Learning Rate)



In [236...

```

class alt7_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(10, kernel_size=3, padding=1)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.LazyConv2d(25, kernel_size=3)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.LazyLinear(200)
        self.fc2 = nn.LazyLinear(120)
        self.fc3 = nn.LazyLinear(84)
        self.fc4 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.relu(self.fc2(out))
        out = self.relu(self.fc3(out))
        out = self.fc4(out)
        return out

model_7 = alt7_LeNet().to(device=try_gpu())
optimizer_7 = optim.SGD(model_7.parameters(), lr=2e-1, weight_decay=1e-4)
model_7.eval()

```

```
Out[236]: alt7_LeNet(
  (conv1): LazyConv2d(0, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): LazyConv2d(0, 25, kernel_size=(3, 3), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): LazyLinear(in_features=0, out_features=200, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc4): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)
```

```
In [237... train_loader_2 = DataLoader(fashion_MNIST, batch_size=32, shuffle=True)
val_loader_2 = DataLoader(fashion_MNIST_val, batch_size=int(len(fashion_MNIST_val)/10)
# Had to reload the loaders with smaller batch sizes because of running low on memory
```

```
In [239... t_loss_hist_7, t_acc_hist_7, v_acc_hist_7 = training_loop(10,
                                                             optimizer_7,
                                                             model_7,
                                                             nn.CrossEntropyLoss(),
                                                             train_loader_2,
                                                             val_loader_2,
                                                             5)
```

Epoch 1:

```
Duration = 13.155 seconds
Training Loss: 0.67564
Training Accuracy: 0.86
Validation Accuracy: 0.854
```

Epoch 5:

```
Duration = 13.621 seconds
Training Loss: 0.24963
Training Accuracy: 0.915
Validation Accuracy: 0.893
```

Epoch 10:


```
Duration = 13.613 seconds
Training Loss: 0.18169
Training Accuracy: 0.937
Validation Accuracy: 0.902
```

Total Training Time = 224.505 seconds

Average Training Time per Epoch = 22.45 seconds

```
In [240... title_7 = "Figure 7 - Loss and Accuracy per Epoch for the altered LeNet Model (Variety
plot_model(title_7, 9, t_loss_hist_7, t_acc_hist_7, v_acc_hist_7, 'right')
```

Figure 7 - Loss and Accuracy per Epoch for the altered LeNet Model (Variety)



Problem 3

In [243...

```

class drop_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(10, kernel_size=3, padding=1)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv1_drop = nn.Dropout2d(p=0.4)
        self.conv2 = nn.LazyConv2d(25, kernel_size=3)
        self.maxp2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2_drop = nn.Dropout2d(p=0.4)
        self.fc1 = nn.LazyLinear(200)
        self.fc1_drop = nn.Dropout(p=0.3)
        self.fc2 = nn.LazyLinear(120)
        self.fc2_drop = nn.Dropout(p=0.2)
        self.fc3 = nn.LazyLinear(84)
        self.fc4 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.conv1_drop(out)
        out = self.maxp2(self.relu(self.conv2(out)))
        out = self.conv2_drop(out)
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.fc1_drop(out)
        out = self.relu(self.fc2(out))
        out = self.fc2_drop(out)

```

```

        out = self.relu(self.fc3(out))
        out = self.fc4(out)
        return out

model_8 = drop_LeNet().to(device=try_gpu())
optimizer_8 = optim.SGD(model_8.parameters(), lr=2e-1, weight_decay=1e-4)
model_8.eval()

```

```

Out[243]: drop_LeNet(
  (conv1): LazyConv2d(0, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv1_drop): Dropout2d(p=0.4, inplace=False)
  (conv2): LazyConv2d(0, 25, kernel_size=(3, 3), stride=(1, 1))
  (maxp2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2_drop): Dropout2d(p=0.4, inplace=False)
  (fc1): LazyLinear(in_features=0, out_features=200, bias=True)
  (fc1_drop): Dropout(p=0.3, inplace=False)
  (fc2): LazyLinear(in_features=0, out_features=120, bias=True)
  (fc2_drop): Dropout(p=0.2, inplace=False)
  (fc3): LazyLinear(in_features=0, out_features=84, bias=True)
  (fc4): LazyLinear(in_features=0, out_features=10, bias=True)
  (relu): ReLU()
  (flat): Flatten(start_dim=1, end_dim=-1)
)

```

```

In [244... t_loss_hist_8, t_acc_hist_8, v_acc_hist_8 = training_loop(10,
                                                             optimizer_8,
                                                             model_8,
                                                             nn.CrossEntropyLoss(),
                                                             train_loader_2,
                                                             val_loader_2,
                                                             5)

```

```

Epoch 1:
  Duration = 13.207 seconds
  Training Loss: 0.67757
  Training Accuracy: 0.871
  Validation Accuracy: 0.861

```

```

Epoch 5:
  Duration = 13.294 seconds
  Training Loss: 0.2448
  Training Accuracy: 0.925
  Validation Accuracy: 0.904

```

```

Epoch 10:
  Duration = 14.133 seconds
  Training Loss: 0.18009
  Training Accuracy: 0.934
  Validation Accuracy: 0.897

```

```

Total Training Time = 228.715 seconds
Average Training Time per Epoch = 22.871 seconds


```

```

In [246... title_8 = "Figure 8 - Loss and Accuracy per Epoch for the altered LeNet Model (Variety
plot_model(title_8, 10, t_loss_hist_8, t_acc_hist_8, v_acc_hist_8, 'right')

```

Figure 8 - Loss and Accuracy per Epoch for the altered LeNet Model (Variety w/ Dropout)



Problem 4

In [279...

```

class drop_LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.LazyConv2d(8, kernel_size=3, stride=2)
        self.maxp1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv1_drop = nn.Dropout2d(p=0.3)
        self.fc1 = nn.LazyLinear(40)
        self.fc2 = nn.LazyLinear(10)
        self.relu = nn.ReLU()
        self.flat = nn.Flatten()

    def forward(self, x):
        out = self.maxp1(self.relu(self.conv1(x)))
        out = self.conv1_drop(out)
        out = self.flat(out)
        out = self.relu(self.fc1(out))
        out = self.fc2(out)
        return out

model_9 = drop_LeNet().to(device=try_gpu())
optimizer_9 = optim.SGD(model_9.parameters(), lr=2e-1)
model_9.eval()

```

```
Out[279]: drop_LeNet(
    (conv1): LazyConv2d(0, 8, kernel_size=(3, 3), stride=(2, 2))
    (maxp1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (conv1_drop): Dropout2d(p=0.3, inplace=False)
    (fc1): LazyLinear(in_features=0, out_features=40, bias=True)
    (fc2): LazyLinear(in_features=0, out_features=10, bias=True)
    (relu): ReLU()
    (flat): Flatten(start_dim=1, end_dim=-1)
)
```

```
In [280... t_loss_hist_9, t_acc_hist_9, v_acc_hist_9 = training_loop(10,
                                                            optimizer_9,
                                                            model_9,
                                                            nn.CrossEntropyLoss(),
                                                            train_loader_1,
                                                            val_loader_1,
                                                            5)
```

Epoch 1:

```
Duration = 7.261 seconds
Training Loss: 0.66293
Training Accuracy: 0.829
Validation Accuracy: 0.819
```

Epoch 5:

```
Duration = 7.352 seconds
Training Loss: 0.36155
Training Accuracy: 0.859
Validation Accuracy: 0.849
```

Epoch 10:


```
Duration = 7.29 seconds
Training Loss: 0.31137
Training Accuracy: 0.872
Validation Accuracy: 0.85
```

Total Training Time = 143.395 seconds

Average Training Time per Epoch = 14.34 seconds

```
In [281... title_9 = "Figure 9 - Loss and Accuracy per Epoch for the simplified LeNet Model"
plot_model(title_9, 11, t_loss_hist_9, t_acc_hist_9, v_acc_hist_9, 'right')
```

Figure 9 - Loss and Accuracy per Epoch for the simplified LeNet Model



```
In [288... with torch.cuda.device(0):
    maccs_LN, params_LN = get_model_complexity_info(model_1, (1, 28, 28), as_strings=True,
                                                    print_per_layer_stat=False, verbose=False)
    maccs_alt, params_alt = get_model_complexity_info(model_7, (1, 28, 28), as_strings=True,
                                                       print_per_layer_stat=False, verbose=False)
    maccs_simp, params_simp = get_model_complexity_info(model_9, (1, 28, 28), as_strings=True,
                                                         print_per_layer_stat=False, verbose=False)
    print('{:<30}  {:<8}'.format('Computational complexity for LeNet: ', maccs_LN))
    print('{:<30}  {:<8}'.format('Number of parameters for LeNet: ', params_LN))
    print('\n{:<30}  {:<8}'.format('Computational complexity for Altered LeNet Model: ', maccs_alt))
    print('{:<30}  {:<8}'.format('Number of parameters for Altered LeNet Model: ', params_alt))
    print('\n{:<30}  {:<8}'.format('Computational complexity for Simple Model: ', maccs_simp))
    print('{:<30}  {:<8}'.format('Number of parameters for Simple Model: ', params_simp))
```

Computational complexity for LeNet: 435.85 KMac
 Number of parameters for LeNet: 61.71 k

Computational complexity for Altered LeNet Model: 644.62 KMac
 Number of parameters for Altered LeNet Model: 217.71 k

Computational complexity for Simple Model: 28.23 KMac
 Number of parameters for Simple Model: 12.05 k