

```
!pip install ptflops
!pip install torchmetrics
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting ptflops
  Downloading ptflops-0.6.9.tar.gz (12 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: torch in /usr/local/lib/python3.9/dist-packages (from ptflops) (1.13.1+cu116)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.9/dist-packages (from torch->ptflops) (4.5.0)
Building wheels for collected packages: ptflops
  Building wheel for ptflops (setup.py) ... done
  Created wheel for ptflops: filename=ptflops-0.6.9-py3-none-any.whl size=11712 sha256=f648b1f4e1604dbb7407f6c832d39776434826d3b0081e3964adbbeb3992d54:
  Stored in directory: /root/.cache/pip/wheels/86/07/9f/879035d99d7b639bbc564d23fed862a679aee7d1a2dced8c2e
Successfully built ptflops
Installing collected packages: ptflops
Successfully installed ptflops-0.6.9
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting torchmetrics
  Downloading torchmetrics-0.11.3-py3-none-any.whl (518 kB)
    

---

 518.6/518.6 KB 4.1 MB/s eta 0:00:00
Requirement already satisfied: torch>=1.8.1 in /usr/local/lib/python3.9/dist-packages (from torchmetrics) (1.13.1+cu116)
Requirement already satisfied: numpy>=1.17.2 in /usr/local/lib/python3.9/dist-packages (from torchmetrics) (1.22.4)
Requirement already satisfied: packaging in /usr/local/lib/python3.9/dist-packages (from torchmetrics) (23.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.9/dist-packages (from torch>=1.8.1->torchmetrics) (4.5.0)
Installing collected packages: torchmetrics
Successfully installed torchmetrics-0.11.3
```

```
import torch
from torch import nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets
import torchvision.transforms as transforms
from torchvision.transforms import ToTensor
import time
from sklearn.model_selection import KFold
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import torch.nn.functional as F
import gc
from ptflops import get_model_complexity_info
from torchmetrics.classification import MulticlassConfusionMatrix
```

▼ Problem 1

```
data_path = '../data-unversioned/ecgr4106/'
cifar10 = datasets.CIFAR10(data_path, train=True, download=True, transform=transforms.Compose([transforms.ToTensor(), transforms.Resize(size=(64, 64))]))
cifar10_val = datasets.CIFAR10(data_path, train=False, download=True, transform=transforms.Compose([transforms.ToTensor(), transforms.Resize(size=(64, 64))]))
```

```
Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to ../data-unversioned/ecgr4106/cif
100% 170498071/170498071 [00:02<00:00, 81632319.27it/s]

Extracting ../data-unversioned/ecgr4106/cifar-10-python.tar.gz to ../data-unversioned/ecgr4106/
Files already downloaded and verified
```

```
def try_gpu(i=0):
    if torch.cuda.device_count() >= i+1:
        return torch.device(f'cuda:{i}')
    return torch.device('cpu')

def training_loop(n_epochs, optimizer, model, loss_fn, train_loader, val_loader, update_freq):
    train_loss_hist = []
    train_acc_hist = []
    val_acc_hist = []
    main_tic = time.perf_counter()

    for epoch in range(1, n_epochs + 1):
        tic = time.perf_counter()
        loss_train = 0.0
        correct_train = 0
        correct_val = 0
        model_argmax = []
        labels_argmax = []

        for imgs, lbls in train_loader:
            images = imgs.to(device=try_gpu())
            labels = lbls.to(device=try_gpu())
            outputs = model(images)
```

```

del images
loss = loss_fn(outputs, labels)
del labels
del outputs
optimizer.zero_grad()
loss.backward()
optimizer.step()
loss_train += loss.item()
gc.collect()
torch.cuda.empty_cache()

toc = time.perf_counter()

with torch.no_grad():
    total = 0
    for imgs, lbls in train_loader:
        images = imgs.to(device=try_gpu())
        labels = lbls.to(device=try_gpu())
        outputs = model(images)
        del images
        _, predicted = torch.max(outputs, dim=1)
        del outputs
        total += labels.shape[0]
        correct_train += int((predicted == labels).sum())
        del labels
        del predicted
    train_acc = round(correct_train/total, 3)
    total = 0
    for imgs, lbls in val_loader:
        images = imgs.to(device=try_gpu())
        labels = lbls.to(device=try_gpu())
        outputs = model(images)
        del images
        _, predicted = torch.max(outputs, dim=1)
        del outputs
        if epoch == 1 or epoch == n_epochs or epoch % update_freq == 0:
            model_argmax = model_argmax + predicted.tolist()
            labels_argmax = labels_argmax + labels.tolist()
        total += labels.shape[0]
        correct_val += int((predicted == labels).sum())
        del labels
        del predicted
    val_acc = round(correct_val/total, 3)

train_loss_hist.append(round(loss_train / len(train_loader), 5))
train_acc_hist.append(train_acc)
val_acc_hist.append(val_acc)
label_set = set(labels_argmax)

if epoch == 1 or epoch == n_epochs or epoch % update_freq == 0:
    print(f"Epoch {epoch}:\n\tDuration = {round(toc - tic, 3)} seconds\n\tTraining Loss: {train_loss_hist[-1]}\n\tTraining Accuracy: {train_acc_hist[-1]}\n\tValidation Loss: {val_loss_hist[-1]}\n\tValidation Accuracy: {val_acc_hist[-1]}\n\tMetric: {metric}")
    print(metric(torch.ByteTensor(model_argmax), torch.ByteTensor(labels_argmax)))

main_toc = time.perf_counter()
print(f"\nTotal Training Time = {round(main_toc - main_tic, 3)} seconds\nAverage Training Time per Epoch (including validation) = {round((main_toc - main_tic) / n_epochs, 3)} seconds")
return train_loss_hist, train_acc_hist, val_acc_hist

def plot_model(title, loss_hist, train_hist, test_hist, leg_loc):
    fig, ax1 = plt.subplots()
    x = range(1, len(loss_hist)+1)
    ax1.plot(x, loss_hist, color='k')
    ax1.set_xlabel('Epoch')
    ax1.set_ylabel('Error')
    ax1.tick_params(axis='y')

    ax2 = ax1.twinx()
    ax2.set_ylabel('Accuracy')
    ax2.plot(x, train_hist)
    ax2.plot(x, test_hist)
    ax2.set_ylim([0, 1])
    ax1.tick_params(axis='y')

    fig.legend(["Training Loss", "Training Accuracy", "Testing Accuracy"], loc=leg_loc, bbox_to_anchor=(1, 1), bbox_transform=ax1.transAxes)
    plt.title(title)

def vgg_block(num_convs, out_channels):
    layers = []
    for _ in range(num_convs):
        layers.append(nn.LazyConv2d(out_channels, kernel_size=3, padding=1))
        layers.append(nn.LazyBatchNorm2d())
        layers.append(nn.ReLU())
    layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
    return nn.Sequential(*layers)

```

```

return nn.Sequential(*layers)

train_loader_1 = DataLoader(cifar10, batch_size=64, shuffle=True)
val_loader_1 = DataLoader(cifar10_val, batch_size=64, shuffle=False)

class tinyVGG(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(tinyVGG, self).__init__()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
        self.fc1 = nn.Linear(128)
        self.fc2 = nn.Linear(64)
        self.fc3 = nn.Linear(num_classes)
        self.fc_drop = nn.Dropout(p=0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.conv_blks(x)
        out = self.fc_drop(self.relu(self.fc1(out)))
        out = self.fc_drop(self.relu(self.fc2(out)))
        out = self.fc3(out)
        return out

model_0 = tinyVGG(arch=((1, 64), (1, 128))).to(device=try_gpu())
optimizer_0 = optim.SGD(model_0.parameters(), lr=0.08)
model_0.eval()

/usr/local/lib/python3.9/dist-packages/torch/nn/modules/lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to
warnings.warn('Lazy modules are a new feature under heavy development '
tinyVGG(
  (conv_blks): Sequential(
    (0): Sequential(
      (0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Flatten(start_dim=1, end_dim=-1)
  )
  (fc1): LazyLinear(in_features=0, out_features=128, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=64, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (fc_drop): Dropout(p=0.5, inplace=False)
  (relu): ReLU()
)

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_0, t_acc_hist_0, v_acc_hist_0= training_loop(10,
                                                    optimizer_0,
                                                    model_0,
                                                    nn.CrossEntropyLoss(),
                                                    train_loader_1,
                                                    val_loader_1,
                                                    1)

title_0 = "Figure 0 - Loss and Accuracy per Epoch for tinyVGG"
plot_model(title_0, t_loss_hist_0, t_acc_hist_0, v_acc_hist_0, 'upper right')

```

```
Epoch 1:
  Duration = 62.585 seconds
  Training Loss: 2.07869
  Training Accuracy: 0.324
  Validation Accuracy: 0.33
tensor([[699, 94, 0, 33, 7, 7, 30, 26, 0, 104],
        [139, 536, 1, 49, 6, 11, 57, 14, 0, 187],
        [243, 122, 21, 135, 98, 38, 242, 55, 0, 46],
        [123, 109, 3, 318, 31, 89, 183, 38, 0, 106],
        [118, 65, 7, 133, 200, 21, 339, 70, 0, 47],
        [128, 88, 6, 251, 47, 236, 148, 33, 0, 63],
        [ 33, 78, 2, 159, 56, 24, 553, 13, 0, 82],
        [132, 117, 1, 118, 109, 30, 82, 236, 0, 175],
        [574, 160, 2, 54, 0, 24, 5, 11, 1, 169],
        [153, 232, 0, 35, 3, 11, 43, 27, 0, 496]])

Epoch 2:
  Duration = 55.115 seconds
  Training Loss: 1.74078
  Training Accuracy: 0.444
  Validation Accuracy: 0.441
tensor([[432, 66, 8, 38, 9, 6, 20, 43, 247, 131],
        [ 15, 588, 7, 28, 2, 11, 11, 24, 70, 244],
        [131, 68, 148, 133, 97, 43, 154, 112, 70, 44],
        [ 47, 53, 31, 376, 22, 100, 151, 66, 63, 91],
        [ 59, 28, 74, 104, 280, 32, 187, 129, 73, 34],
        [ 35, 43, 41, 250, 27, 252, 97, 118, 91, 46],
        [ 7, 38, 38, 135, 60, 32, 563, 29, 39, 59],
        [ 48, 48, 16, 93, 47, 35, 34, 519, 35, 125],
        [ 55, 79, 0, 42, 3, 7, 8, 12, 657, 137],
        [ 14, 235, 2, 22, 5, 7, 16, 26, 78, 595]])

Epoch 3:
  Duration = 53.337 seconds
  Training Loss: 1.53197
  Training Accuracy: 0.495
  Validation Accuracy: 0.493
tensor([[586, 55, 56, 38, 6, 7, 4, 26, 177, 45],
        [ 44, 633, 6, 25, 2, 10, 3, 15, 61, 201],
        [ 80, 39, 390, 159, 63, 72, 49, 90, 31, 27],
        [ 45, 28, 82, 557, 19, 128, 24, 47, 31, 39],
        [ 62, 17, 167, 145, 292, 69, 57, 139, 37, 15],
        [ 32, 23, 78, 341, 26, 348, 9, 90, 30, 23],
        [ 12, 31, 85, 313, 91, 73, 322, 32, 12, 29],
        [ 56, 30, 49, 145, 26, 61, 2, 558, 21, 52],
        [107, 73, 11, 44, 6, 8, 1, 11, 685, 54],
        [ 55, 221, 14, 33, 4, 3, 2, 30, 80, 558]])

Epoch 4:
  Duration = 54.212 seconds
  Training Loss: 1.39665
  Training Accuracy: 0.549
  Validation Accuracy: 0.536
tensor([[548, 8, 83, 22, 55, 28, 34, 43, 138, 41],
        [ 47, 554, 30, 17, 13, 25, 24, 45, 54, 191],
        [ 42, 3, 352, 38, 253, 121, 109, 63, 12, 7],
        [ 11, 2, 75, 246, 112, 350, 137, 52, 6, 9],
        [ 28, 1, 93, 25, 573, 77, 125, 63, 14, 1],
        [ 1, 2, 77, 81, 119, 590, 57, 59, 8, 6],
        [ 3, 3, 58, 44, 146, 74, 650, 8, 6, 8],
        [ 12, 2, 50, 30, 134, 123, 28, 601, 6, 14],
        [114, 33, 18, 19, 36, 42, 22, 20, 637, 59],
        [ 41, 111, 20, 30, 17, 25, 35, 63, 52, 606]])

Epoch 5:
  Duration = 53.906 seconds
  Training Loss: 1.26966
  Training Accuracy: 0.592
  Validation Accuracy: 0.558
tensor([[534, 13, 134, 47, 25, 19, 50, 11, 123, 44],
        [ 34, 593, 15, 28, 6, 2, 35, 14, 52, 221],
        [ 38, 3, 533, 94, 84, 71, 123, 30, 14, 10],
        [ 11, 6, 96, 487, 54, 159, 148, 14, 9, 16],
        [ 23, 4, 236, 76, 384, 50, 151, 56, 14, 6],
        [ 8, 1, 104, 248, 53, 447, 80, 35, 14, 10],
        [ 7, 5, 71, 81, 53, 31, 737, 5, 4, 6],
        [ 12, 2, 74, 93, 84, 96, 65, 537, 9, 28],
        [ 90, 51, 26, 41, 14, 15, 23, 7, 670, 63],
        [ 33, 97, 24, 40, 5, 7, 50, 37, 49, 658]])

Epoch 6:
  Duration = 53.128 seconds
  Training Loss: 1.16226
  Training Accuracy: 0.635
  Validation Accuracy: 0.591
tensor([[660, 11, 63, 20, 17, 5, 24, 12, 167, 21],
        [ 36, 706, 16, 23, 7, 1, 10, 15, 99, 87],
        [ 68, 9, 470, 63, 177, 36, 91, 49, 31, 6],
        [ 22, 8, 104, 485, 97, 73, 127, 46, 25, 13],
        [ 33, 3, 115, 61, 580, 11, 99, 75, 22, 1],
        [ 15, 6, 104, 296, 94, 306, 72, 73, 26, 8],
        [ 9, 4, 63, 73, 115, 9, 692, 14, 14, 7],
        [ 20, 3, 53, 75, 102, 34, 33, 654, 12, 14],
        [ 82, 33, 14, 14, 16, 3, 9, 4, 808, 17],
        [ 58, 155, 21, 27, 10, 0, 17, 41, 120, 551]])

Epoch 7:
```

```

Duration = 53.04 seconds
Training Loss: 1.07017
Training Accuracy: 0.671
Validation Accuracy: 0.602
tensor([[766, 29, 48, 32, 14, 7, 24, 18, 31, 31],
        [ 50, 772, 19, 21, 4, 5, 7, 11, 8, 103],
        [ 86, 8, 565, 79, 72, 46, 92, 37, 5, 10],
        [ 30, 9, 119, 450, 64, 163, 110, 34, 3, 18],
        [ 43, 5, 208, 70, 450, 25, 115, 77, 5, 2],
        [ 13, 7, 99, 229, 55, 467, 58, 62, 3, 7],
        [ 12, 8, 91, 88, 53, 20, 707, 12, 2, 7],
        [ 27, 1, 64, 77, 53, 57, 24, 669, 1, 27],
        [215, 80, 28, 43, 17, 15, 13, 4, 546, 39],
        [ 66, 179, 18, 27, 8, 10, 12, 36, 19, 625]])

```

```

Epoch 8:
Duration = 54.059 seconds
Training Loss: 0.97531
Training Accuracy: 0.689
Validation Accuracy: 0.608
tensor([[741, 13, 47, 10, 1, 5, 18, 5, 135, 25],
        [ 52, 673, 15, 12, 2, 7, 3, 1, 82, 153],
        [ 93, 8, 618, 66, 34, 44, 65, 22, 35, 15],
        [ 41, 12, 139, 451, 30, 131, 100, 19, 50, 27],
        [ 57, 7, 266, 80, 360, 28, 118, 51, 29, 4],
        [ 34, 7, 114, 232, 34, 454, 58, 25, 29, 13],
        [ 13, 11, 111, 83, 21, 19, 712, 5, 13, 12],
        [ 46, 4, 88, 76, 62, 69, 29, 568, 16, 42],
        [ 92, 33, 9, 3, 2, 4, 6, 2, 819, 30],
        [ 76, 99, 15, 9, 2, 6, 14, 13, 87, 679]])

```

```

Epoch 9:
Duration = 53.973 seconds
Training Loss: 0.89546
Training Accuracy: 0.742
Validation Accuracy: 0.635
tensor([[685, 19, 47, 12, 9, 4, 22, 9, 164, 29],
        [ 39, 758, 11, 11, 3, 4, 6, 4, 71, 93],
        [ 79, 10, 523, 43, 127, 40, 97, 36, 35, 10],
        [ 26, 14, 99, 420, 87, 126, 138, 39, 36, 15],
        [ 39, 6, 108, 49, 588, 12, 111, 62, 22, 3],
        [ 26, 4, 92, 189, 89, 430, 77, 56, 26, 11],
        [ 9, 13, 54, 47, 68, 10, 770, 8, 11, 10],
        [ 29, 4, 51, 41, 74, 43, 27, 693, 18, 20],
        [ 60, 37, 8, 7, 6, 5, 14, 3, 838, 22],
        [ 38, 150, 10, 15, 3, 5, 18, 26, 90, 645]])

```

```

Epoch 10:
Duration = 53.07 seconds
Training Loss: 0.8161
Training Accuracy: 0.774
Validation Accuracy: 0.65
tensor([[752, 15, 40, 19, 24, 8, 29, 22, 64, 27],
        [ 46, 738, 10, 16, 10, 7, 14, 10, 41, 108],
        [ 80, 5, 446, 66, 177, 59, 97, 53, 10, 7],
        [ 18, 5, 58, 484, 96, 134, 130, 47, 14, 14],
        [ 35, 2, 55, 55, 652, 21, 91, 82, 6, 1],
        [ 13, 2, 46, 202, 87, 500, 74, 66, 5, 5],
        [ 6, 5, 27, 55, 105, 16, 761, 15, 6, 4],
        [ 20, 2, 26, 37, 77, 53, 24, 746, 4, 11],
        [102, 41, 13, 22, 14, 11, 23, 5, 742, 27],
        [ 59, 116, 12, 14, 8, 6, 25, 33, 46, 681]])

```

Total Training Time: 303.770 seconds

```

class VGG_11(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(VGG_11, self).__init__()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
        self.fc1 = nn.Linear(4096)
        self.fc2 = nn.Linear(4096)
        self.fc3 = nn.Linear(num_classes)
        self.fc_drop = nn.Dropout(p=0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.conv_blks(x)
        out = self.fc_drop(self.relu(self.fc1(out)))
        out = self.fc_drop(self.relu(self.fc2(out)))
        out = self.fc3(out)
        return out

model_1 = VGG_11(arch=((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))).to(device=try_gpu())
optimizer_1 = optim.SGD(model_1.parameters(), lr=0.1)
model_1.eval()

```

```

VGG_11(
  (conv_blks): Sequential(
    (0): Sequential(

```

```

(0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(2): ReLU()
(3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(1): Sequential(
  (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(2): Sequential(
  (0): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (5): ReLU()
  (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(3): Sequential(
  (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (5): ReLU()
  (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(4): Sequential(
  (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (5): ReLU()
  (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(5): Flatten(start_dim=1, end_dim=-1)
)
(fc1): LazyLinear(in_features=0, out_features=4096, bias=True)
(fc2): LazyLinear(in_features=0, out_features=4096, bias=True)
(fc3): LazyLinear(in_features=0, out_features=10, bias=True)
(fc_drop): Dropout(p=0.5, inplace=False)
(relu): ReLU()
)

```

```

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_1, t_acc_hist_1, v_acc_hist_1 = training_loop(10,
                                                         optimizer_1,
                                                         model_1,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         2)

title_1 = "Figure 1 - Loss and Accuracy per Epoch for VGG-11"
plot_model(title_1, t_loss_hist_1, t_acc_hist_1, v_acc_hist_1, 'upper right')

```

```
Epoch 1:
Duration = 78.519 seconds
Training Loss: 2.30272
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0],
[ 0, 1000,  0,  0,  0,  0,  0,  0,  0,  0]])
```

```
Epoch 2:
Duration = 78.586 seconds
Training Loss: 2.30284
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0]])
```

```
Epoch 4:
Duration = 78.412 seconds
Training Loss: 2.30283
Training Accuracy: 0.103
Validation Accuracy: 0.103
tensor([[ 0,  0,  0,  0,  0,  984,  0,  0,  0,  16],
[ 0,  0,  0,  0,  0,  978,  0,  0,  0,  22],
[ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
[ 0,  0,  0,  0,  0,  998,  0,  0,  0,  2],
[ 0,  0,  0,  0,  0,  999,  0,  0,  0,  1],
[ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
[ 0,  0,  0,  0,  0,  999,  0,  0,  0,  1],
[ 0,  0,  0,  0,  0,  990,  0,  0,  0,  10],
[ 0,  0,  0,  0,  0,  963,  0,  0,  0,  37],
[ 0,  0,  0,  0,  0,  970,  0,  0,  0,  30]])
```

```
Epoch 6:
Duration = 76.961 seconds
Training Loss: 2.30278
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0]])
```

```
Epoch 8:
Duration = 77.631 seconds
Training Loss: 2.30272
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0]])
```

```
Epoch 10:
Duration = 77.327 seconds
Training Loss: 2.30252
Training Accuracy: 0.112
Validation Accuracy: 0.111
```

```
class VGG_16(nn.Module):
def __init__(self, arch, num_classes=10):
super(VGG_16, self).__init__()
conv_blks = []
for (num_convs, out_channels) in arch:
conv_blks.append(vgg_block(num_convs, out_channels))
self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
self.fc1 = nn.LazyLinear(4096)
self.fc2 = nn.LazyLinear(4096)
```

```

self.fc3 = nn.LazyLinear(num_classes)
self.fc_drop = nn.Dropout(p=0.5)
self.relu = nn.ReLU()

def forward(self, x):
    out = self.conv_blks(x)
    out = self.fc_drop(self.relu(self.fc1(out)))
    out = self.fc_drop(self.relu(self.fc2(out)))
    out = self.fc3(out)
    return out

model_2 = VGG_16(arch=((2, 64), (2, 128), (3, 256), (3, 512), (3, 512))).to(device=try_gpu())
optimizer_2 = optim.SGD(model_2.parameters(), lr=0.1)
model_2.eval()

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_2, t_acc_hist_2, v_acc_hist_2 = training_loop(3,
                                                         optimizer_2,
                                                         model_2,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)

title_2 = "Figure 2 - Loss and Accuracy per Epoch for VGG-16"
plot_model(title_2, t_loss_hist_2, t_acc_hist_2, v_acc_hist_2, 'upper right')

```

```

class VGG_19(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(VGG_19, self).__init__()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
        self.fc1 = nn.LazyLinear(4096)
        self.fc2 = nn.LazyLinear(4096)
        self.fc3 = nn.LazyLinear(num_classes)
        self.fc_drop = nn.Dropout(p=0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.conv_blks(x)
        out = self.fc_drop(self.relu(self.fc1(out)))
        out = self.fc_drop(self.relu(self.fc2(out)))
        out = self.fc3(out)
        return out

model_3 = VGG_19(arch=((2, 64), (2, 128), (4, 256), (4, 512), (4, 512))).to(device=try_gpu())
optimizer_3 = optim.SGD(model_3.parameters(), lr=0.1)
model_3.eval()

```

```

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_3, t_acc_hist_3, v_acc_hist_3 = training_loop(3,
                                                         optimizer_3,
                                                         model_3,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)

title_3 = "Figure 3 - Loss and Accuracy per Epoch for VGG-19"
plot_model(title_3, t_loss_hist_3, t_acc_hist_3, v_acc_hist_3, 'upper right')

```

▼ Problem 2

```

class Inception(nn.Module):
    def __init__(self, c1, c2, c3, c4, **kwargs):
        super(Inception, self).__init__(**kwargs)
        self.b1_1 = nn.LazyConv2d(c1, kernel_size=1)
        self.b2_1 = nn.LazyConv2d(c2[0], kernel_size=1)
        self.b2_2 = nn.LazyConv2d(c2[1], kernel_size=3, padding=1)
        self.b3_1 = nn.LazyConv2d(c3[0], kernel_size=1)
        self.b3_2 = nn.LazyConv2d(c3[1], kernel_size=5, padding=2)
        self.b4_1 = nn.MaxPool2d(kernel_size=3, stride=1, padding=1)
        self.b4_2 = nn.LazyConv2d(c4, kernel_size=1)

    def forward(self, x):
        b1 = F.relu(self.b1_1(x))
        b2 = F.relu(self.b2_2(F.relu(self.b2_1(x))))

```



```

b3 = F.relu(self.b3_2(F.relu(self.b3_1(x))))
b4 = F.relu(self.b4_2(self.b4_1(x)))
return torch.cat((b1, b2, b3, b4), dim=1)

```

```

class GoogLeNet(nn.Module):
    def __init__(self, num_classes=10):
        super(GoogLeNet, self).__init__()
        self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
                                   nn.LazyConv2d(64, kernel_size=1),
                                   nn.ReLU(),
                                   nn.LazyConv2d(192, kernel_size=3, padding=1),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body1 = nn.Sequential(Inception(64, (96, 128), (16, 32), 32),
                                    Inception(128, (128, 192), (32, 96), 64),
                                    nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body2 = nn.Sequential(Inception(192, (96, 208), (16, 48), 64),
                                    Inception(160, (112, 124), (24, 64), 64),
                                    Inception(128, (128, 256), (24, 64), 64),
                                    Inception(112, (144, 288), (32, 64), 64),
                                    Inception(256, (160, 320), (32, 128), 128),
                                    nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body3 = nn.Sequential(Inception(256, (160, 320), (32, 128), 128),
                                    Inception(384, (192, 384), (48, 128), 128),
                                    nn.AdaptiveAvgPool2d((1,1)),
                                    nn.Flatten())
        self.fc = nn.LazyLinear(num_classes)

    def forward(self, x):
        out = self.stem(x)
        out = self.body1(out)
        out = self.body2(out)
        out = self.body3(out)
        out = self.fc(out)
        return out

```

```

model_4 = GoogLeNet().to(device=try_gpu())
optimizer_4 = optim.SGD(model_4.parameters(), lr=0.1)
model_4.eval()

```

```

    )
    (1): Inception(
      (b1_1): LazyConv2d(0, 384, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 192, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 48, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
    )
    (2): AdaptiveAvgPool2d(output_size=(1, 1))
    (3): Flatten(start_dim=1, end_dim=-1)
  )
  (fc): LazyLinear(in_features=0, out_features=10, bias=True)
)

```

```
torch.cuda.empty_cache()
```

```
gc.collect()
```

```

t_loss_hist_4, t_acc_hist_4, v_acc_hist_4 = training_loop(3,
                                                         optimizer_4,
                                                         model_4,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)

```

```
title_4 = "Figure 4 - Loss and Accuracy per Epoch for GoogLeNet"
```

```
plot_model(title_4, t_loss_hist_4, t_acc_hist_4, v_acc_hist_4, 'upper right')
```

```

Epoch 1:
    Duration = 60.74 seconds
    Training loss: 2.30782

class AltGoogLeNet(nn.Module):
    def __init__(self, num_classes=10):
        super(GoogLeNet, self).__init__()
        self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
                                   nn.LazyConv2d(64, kernel_size=1),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.LazyConv2d(192, kernel_size=3, padding=1),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body1 = nn.Sequential(Inception(64, (96, 128), (16, 32), 32),
                                   Inception(128, (128, 192), (32, 96), 64),
                                   nn.LazyBatchNorm2d(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body2 = nn.Sequential(Inception(192, (96, 208), (16, 48), 64),
                                   Inception(160, (112, 124), (24, 64), 64),
                                   Inception(128, (128, 256), (24, 64), 64),
                                   Inception(112, (144, 288), (32, 64), 64),
                                   Inception(256, (160, 320), (32, 128), 128),
                                   nn.LazyBatchNorm2d(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body3 = nn.Sequential(Inception(256, (160, 320), (32, 128), 128),
                                   Inception(384, (192, 384), (48, 128), 128),
                                   nn.LazyBatchNorm2d(),
                                   nn.AdaptiveAvgPool2d((1,1)),
                                   nn.Flatten())
        self.fc = nn.LazyLinear(num_classes)

    def forward(self, x):
        out = self.stem(x)
        out = self.body1(out)
        out = self.body2(out)
        out = self.body3(out)
        out = self.fc(out)
        return out

model_5 = GoogLeNet().to(device=try_gpu())
optimizer_5 = optim.SGD(model_5.parameters(), lr=0.1)
model_5.eval()

```

```

(0): Inception(
  (b1_1): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(1, 1))
  (b2_1): LazyConv2d(0, 160, kernel_size=(1, 1), stride=(1, 1))
  (b2_2): LazyConv2d(0, 320, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (b3_1): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
  (b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
  (b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
)
(1): Inception(
  (b1_1): LazyConv2d(0, 384, kernel_size=(1, 1), stride=(1, 1))
  (b2_1): LazyConv2d(0, 192, kernel_size=(1, 1), stride=(1, 1))
  (b2_2): LazyConv2d(0, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (b3_1): LazyConv2d(0, 48, kernel_size=(1, 1), stride=(1, 1))
  (b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
  (b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
)
(2): AdaptiveAvgPool2d(output_size=(1, 1))
(3): Flatten(start dim=1, end dim=-1)

```

```

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_5, t_acc_hist_5, v_acc_hist_5 = training_loop(10,
    optimizer_5,
    model_5,
    nn.CrossEntropyLoss(),
    train_loader_1,
    val_loader_1,
    2)
title_5 = "Figure 5 - Loss and Accuracy per Epoch for GoogLeNet with Batch Norm"
plot_model(title_5, t_loss_hist_5, t_acc_hist_5, v_acc_hist_5, 'upper right')

```

```
Epoch 1:
Duration = 59.662 seconds
Training Loss: 2.30278
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0]])
```

```
Epoch 2:
Duration = 59.788 seconds
Training Loss: 2.3028
Training Accuracy: 0.1
Validation Accuracy: 0.1
tensor([[ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0],
        [ 0,  0,  0,  0,  0,  0,  0, 1000,  0,  0]])
```

▼ Problem 3

```

class Residual(nn.Module):
    def __init__(self, num_channels, use_1x1conv=False, strides=1):
        super().__init__()
        self.conv1 = nn.Conv2d(num_channels, kernel_size=3, padding=1, stride=strides)
        self.conv2 = nn.Conv2d(num_channels, kernel_size=3, padding=1)
        if use_1x1conv:
            self.conv3 = nn.Conv2d(num_channels, kernel_size=1, stride=strides)
        else:
            self.conv3 = None
        self.bn = nn.BatchNorm2d(num_channels)

    def forward(self, x):
        y = F.relu(self.bn(self.conv1(x)))
        y = self.bn(self.conv2(y))
        if self.conv3:
            x = self.bn(self.conv3(x))
        y += x
        return F.relu(y)

```

```

def block(num_residuals, num_channels, first_block=False):
    blk = []
    for i in range(num_residuals):
        if i == 0 and not first_block:
            blk.append(Residual(num_channels, use_1x1conv=True, strides=2))
        else:
            blk.append(Residual(num_channels))
    return nn.Sequential(*blk)

```

```

class ResNet(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(ResNet, self).__init__()
        self.stem = nn.Sequential(nn.Conv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.BatchNorm2d(64),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

        blks = []
        for i, b in enumerate(arch):
            blks.append(block(*b, first_block=(i==0)))
        self.blks = nn.Sequential(*blks)
        self.head = nn.Sequential(nn.AdaptiveAvgPool2d((1, 1)),
                                   nn.Flatten(),
                                   nn.Linear(num_classes))

    def forward(self, x):
        out = self.stem(x)
        out = self.blks(out)
        out = self.head(out)
        return out

```

```

model_6 = ResNet(arch=((2, 64), (2, 128), (2, 256), (2, 512))).to(device=try_gpu())
optimizer_6 = optim.SGD(model_6.parameters(), lr=0.01)
model_6.eval()

```

```

        (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): Residual(
      (conv1): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Sequential(
    (0): Residual(
      (conv1): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
      (conv2): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv3): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(2, 2))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): Residual(
      (conv1): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv2): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (2): Sequential(
    (0): Residual(
      (conv1): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
      (conv2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv3): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(2, 2))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): Residual(
      (conv1): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (3): Sequential(
    (0): Residual(
      (conv1): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
      (conv2): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv3): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(2, 2))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): Residual(
      (conv1): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv2): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  )
  (head): Sequential(
    (0): AdaptiveAvgPool2d(output_size=(1, 1))
    (1): Flatten(start_dim=1, end_dim=-1)
    (2): LazyLinear(in_features=0, out_features=10, bias=True)
  )
)

```

```

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_6, t_acc_hist_6, v_acc_hist_6 = training_loop(10,
                                                         optimizer_6,
                                                         model_6,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         2)
title_6 = "Figure 6 - Loss and Accuracy per Epoch for ResNet"
plot_model(title_6, t_loss_hist_6, t_acc_hist_6, v_acc_hist_6, 'upper right')

```

```
Epoch 1:
Duration = 59.61 seconds
Training Loss: 2.30173
Training Accuracy: 0.149
Validation Accuracy: 0.153
tensor([[637, 0, 0, 0, 363, 0, 0, 0, 0, 0],
        [325, 0, 0, 0, 675, 0, 0, 0, 0, 0],
        [232, 0, 0, 0, 768, 0, 0, 0, 0, 0],
        [177, 0, 0, 0, 823, 0, 0, 0, 0, 0],
        [109, 0, 0, 0, 891, 0, 0, 0, 0, 0],
        [164, 0, 0, 0, 836, 0, 0, 0, 0, 0],
        [112, 0, 0, 0, 888, 0, 0, 0, 0, 0],
        [268, 0, 0, 0, 732, 0, 0, 0, 0, 0],
        [597, 0, 0, 0, 403, 0, 0, 0, 0, 0],
        [522, 0, 0, 0, 478, 0, 0, 0, 0, 0]])
```

```
Epoch 2:
Duration = 57.799 seconds
Training Loss: 2.29885
Training Accuracy: 0.126
Validation Accuracy: 0.129
tensor([[972, 0, 0, 0, 28, 0, 0, 0, 0, 0],
        [877, 0, 0, 0, 121, 0, 0, 0, 0, 2],
        [797, 0, 0, 0, 203, 0, 0, 0, 0, 0],
        [762, 0, 0, 0, 237, 0, 0, 0, 0, 1],
        [684, 0, 0, 0, 315, 0, 0, 0, 0, 1],
        [831, 0, 0, 0, 167, 0, 0, 0, 0, 2],
        [659, 0, 0, 0, 341, 0, 0, 0, 0, 0],
        [884, 0, 0, 0, 116, 0, 0, 0, 0, 0],
        [962, 0, 0, 0, 38, 0, 0, 0, 0, 0],
        [955, 0, 0, 0, 45, 0, 0, 0, 0, 0]])
```

```
Epoch 4:
Duration = 59.47 seconds
Training Loss: 2.28171
Training Accuracy: 0.154
Validation Accuracy: 0.156
tensor([[924, 0, 0, 0, 7, 3, 38, 0, 1, 27],
        [734, 0, 0, 0, 33, 12, 145, 0, 12, 64],
        [555, 0, 0, 0, 85, 14, 237, 0, 2, 107],
        [555, 0, 0, 1, 64, 34, 241, 0, 5, 100],
        [416, 0, 0, 1, 104, 16, 344, 0, 0, 119],
        [620, 0, 0, 1, 50, 48, 186, 0, 2, 93],
        [359, 0, 0, 0, 100, 16, 404, 0, 1, 120],
        [687, 0, 0, 0, 68, 13, 132, 0, 2, 98],
        [939, 0, 0, 0, 5, 5, 36, 0, 5, 10],
        [842, 0, 0, 0, 27, 7, 46, 0, 4, 74]])
```

```
Epoch 6:
Duration = 58.969 seconds
Training Loss: 2.12357
Training Accuracy: 0.208
Validation Accuracy: 0.208
tensor([[748, 36, 1, 0, 0, 27, 39, 15, 90, 44],
        [363, 104, 3, 0, 16, 91, 115, 30, 191, 87],
        [239, 56, 13, 0, 11, 149, 380, 21, 56, 75],
        [167, 70, 8, 0, 35, 207, 300, 45, 54, 114],
        [117, 37, 10, 0, 18, 116, 560, 26, 41, 75],
        [153, 125, 9, 0, 18, 202, 288, 29, 97, 79],
        [63, 24, 1, 0, 26, 117, 635, 38, 17, 79],
        [142, 66, 4, 0, 22, 114, 219, 78, 119, 236],
        [729, 35, 2, 0, 3, 26, 16, 16, 116, 57],
        [415, 33, 4, 0, 13, 42, 71, 39, 214, 169]])
```

```
Epoch 8:
Duration = 58.231 seconds
Training Loss: 2.03023
Training Accuracy: 0.25
Validation Accuracy: 0.254
tensor([[49, 204, 14, 22, 0, 73, 56, 46, 332, 204],
        [1, 377, 0, 6, 0, 169, 138, 106, 90, 113],
        [21, 67, 19, 37, 0, 148, 506, 92, 56, 54],
        [3, 63, 5, 49, 0, 229, 455, 149, 7, 40],
        [2, 36, 3, 23, 0, 94, 681, 90, 26, 45],
        [4, 69, 3, 32, 0, 310, 443, 102, 10, 27],
        [1, 10, 0, 33, 0, 67, 766, 101, 4, 18],
        [2, 60, 3, 20, 0, 130, 312, 317, 12, 144],
        [7, 238, 6, 13, 0, 85, 27, 81, 296, 247],
        [0, 174, 0, 15, 0, 69, 104, 202, 84, 352]])
```

```
Epoch 10:
Duration = 58.95 seconds
Training Loss: 1.95409
Training Accuracy: 0.289
Validation Accuracy: 0.295
tensor([[400, 142, 6, 15, 1, 94, 44, 35, 127, 136],
        [32, 445, 3, 6, 0, 136, 94, 79, 103, 102],
        [103, 64, 13, 22, 0, 238, 426, 74, 16, 44],
        [26, 53, 8, 32, 1, 362, 351, 117, 8, 42],
        [28, 29, 4, 11, 0, 163, 625, 88, 18, 34],
        [19, 72, 6, 14, 0, 422, 347, 81, 11, 28],
        [9, 19, 3, 21, 0, 149, 694, 88, 1, 16],
        [18, 89, 3, 13, 0, 174, 254, 322, 11, 116]])
```

```
class Residual26(nn.Module):
    def __init__(self, in_channels, out_channels, use_1x1conv=False, strides=1):
```



```

        nn.ReLU(),
        nn.LazyConv2d(256, kernel_size=3, padding=1),
        nn.LazyBatchNorm2d(),
        nn.ReLU(),
        nn.LazyConv2d(1024, kernel_size=1),
        nn.LazyBatchNorm2d())
self.b7_1 = nn.Sequential(nn.LazyConv2d(512, kernel_size=1),
        nn.LazyBatchNorm2d(),
        nn.ReLU(),
        nn.LazyConv2d(512, kernel_size=3, padding=1),
        nn.LazyBatchNorm2d(),
        nn.ReLU(),
        nn.LazyConv2d(2048, kernel_size=1),
        nn.LazyBatchNorm2d())
self.b7_2 = nn.Sequential(nn.LazyConv2d(2048, kernel_size=1),
        nn.LazyBatchNorm2d())
self.b8 = nn.Sequential(nn.LazyConv2d(512, kernel_size=1),
        nn.LazyBatchNorm2d(),
        nn.ReLU(),
        nn.LazyConv2d(512, kernel_size=3, padding=1),
        nn.LazyBatchNorm2d(),
        nn.ReLU(),
        nn.LazyConv2d(2048, kernel_size=1),
        nn.LazyBatchNorm2d())

self.relu = nn.ReLU()
self.head = nn.Sequential(nn.AdaptiveAvgPool2d((1, 1)),
        nn.Flatten(),
        nn.LazyLinear(num_classes))

def forward(self, x):
    out = self.stem(x)
    out = self.relu(self.b1_1(out) + self.b1_2(out))
    out = self.relu(out + self.b2(out))
    out = self.relu(self.b3_1(out) + self.b3_2(out))
    out = self.relu(out + self.b4(out))
    out = self.relu(self.b5_1(out) + self.b5_2(out))
    out = self.relu(out + self.b6(out))
    out = self.relu(self.b7_1(out) + self.b7_2(out))
    out = self.relu(out + self.b8(out))
    out = self.head(out)
    return out

model_7 = ResNet26().to(device=try_gpu())
optimizer_7 = optim.SGD(model_7.parameters(), lr=0.1)
model_7.eval()

(3): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(5): ReLU()
(6): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(1, 1))
(7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

(6): LazyConv2d(0, 1024, kernel_size=(1, 1), stride=(1, 1))
(7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(b7_1): Sequential(
  (0): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (5): ReLU()
  (6): LazyConv2d(0, 2048, kernel_size=(1, 1), stride=(1, 1))
  (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(b7_2): Sequential(
  (0): LazyConv2d(0, 2048, kernel_size=(1, 1), stride=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)

torch.cuda.empty_cache()
gc.collect()
t_loss_hist_7, t_acc_hist_7, v_acc_hist_7 = training_loop(5,
                                                         optimizer_7,
                                                         model_7,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         3)

title_7 = "Figure 7 - Loss and Accuracy per Epoch for ResNet26"
plot_model(title_7, t_loss_hist_7, t_acc_hist_7, v_acc_hist_7, 'upper right')

```

```
Epoch 1:
  Duration = 317.779 seconds
  Training Loss: 2.15581
  Training Accuracy: 0.194
  Validation Accuracy: 0.194
```

```
class ResNet34(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(ResNet34, self).__init__()
        self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

        blks = []
        for i, b in enumerate(arch):
            blks.append(block(*b, first_block=(i==0)))
        self.blks = nn.Sequential(*blks)
        self.head = nn.Sequential(nn.AdaptiveAvgPool2d((1, 1)),
                                   nn.Flatten(),
                                   nn.LazyLinear(num_classes))

    def forward(self, x):
        out = self.stem(x)
        out = self.blks(out)
        out = self.head(out)
        return out

model_8 = ResNet34(arch=((3, 64), (4, 128), (6, 256), (3, 512))).to(device=try_gpu())
optimizer_8 = optim.SGD(model_8.parameters(), lr=0.01)
model_8.eval()

    (0): Residual(
      (conv1): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
      (conv2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv3): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(2, 2))
      (bn): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
```

```
torch.cuda.empty_cache()
gc.collect()
t_loss_hist_8, t_acc_hist_8, v_acc_hist_8 = training_loop(10,
                                                           optimizer_8,
                                                           model_8,
                                                           nn.CrossEntropyLoss(),
                                                           train_loader_1,
                                                           val_loader_1,
                                                           2)

title_8 = "Figure 8 - Loss and Accuracy per Epoch for ResNet34"
plot_model(title_8, t_loss_hist_8, t_acc_hist_8, v_acc_hist_8, 'upper right')
```



Epoch 1:
Duration = 66.322 seconds
Training Loss: 1.85303
Training Accuracy: 0.33
Validation Accuracy: 0.338
tensor([[521, 122, 9, 27, 1, 23, 9, 30, 204, 54],
[104, 512, 3, 24, 1, 31, 8, 13, 180, 124],
[139, 82, 55, 103, 90, 185, 152, 112, 47, 35],
[86, 68, 25, 260, 12, 260, 90, 108, 47, 44],
[65, 35, 19, 101, 162, 163, 243, 132, 49, 31],
[61, 79, 25, 178, 16, 347, 81, 116, 53, 44],
[27, 57, 9, 140, 65, 199, 355, 117, 15, 16],
[69, 96, 14, 94, 21, 125, 48, 360, 31, 142],
[279, 122, 1, 39, 0, 18, 0, 16, 461, 64],
[99, 185, 12, 31, 1, 29, 4, 62, 226, 351]])

Epoch 2:
Duration = 65.853 seconds
Training Loss: 1.81662
Training Accuracy: 0.345
Validation Accuracy: 0.352
tensor([[512, 79, 17, 58, 16, 26, 8, 77, 124, 83],
[116, 435, 3, 47, 4, 35, 14, 52, 83, 211],
[90, 44, 42, 150, 275, 112, 95, 161, 15, 16],
[28, 25, 21, 308, 102, 165, 140, 167, 14, 30],
[49, 11, 11, 102, 421, 77, 132, 162, 19, 16],
[23, 28, 11, 245, 124, 268, 102, 169, 15, 15],
[5, 14, 6, 150, 316, 73, 279, 142, 2, 13],
[29, 36, 9, 108, 101, 75, 64, 517, 9, 52],
[298, 85, 10, 91, 1, 24, 5, 34, 325, 127],
[86, 142, 4, 59, 4, 25, 14, 168, 84, 414]])

Epoch 4:
Duration = 66.722 seconds
Training Loss: 1.7389
Training Accuracy: 0.388
Validation Accuracy: 0.393
tensor([[420, 102, 30, 28, 4, 18, 18, 44, 241, 95],
[38, 561, 0, 20, 0, 17, 24, 11, 78, 251],
[72, 56, 114, 118, 168, 102, 191, 104, 40, 35],
[30, 52, 38, 303, 34, 169, 208, 91, 27, 48],
[39, 23, 45, 102, 241, 77, 280, 125, 40, 28],
[13, 42, 39, 226, 40, 268, 184, 108, 38, 42],
[7, 39, 19, 153, 85, 57, 532, 73, 6, 29],
[23, 43, 35, 107, 37, 85, 112, 428, 17, 113],
[166, 110, 14, 41, 0, 18, 7, 15, 491, 138],
[29, 195, 4, 33, 0, 15, 23, 47, 82, 572]])

Epoch 6:
Duration = 67.3 seconds
Training Loss: 1.66135
Training Accuracy: 0.397
Validation Accuracy: 0.393
tensor([[415, 28, 112, 61, 8, 11, 33, 76, 163, 93],
[41, 381, 13, 58, 0, 22, 70, 50, 57, 308],
[37, 6, 230, 127, 205, 54, 186, 118, 19, 18],
[8, 10, 89, 302, 70, 86, 284, 123, 8, 20],
[31, 3, 94, 89, 336, 36, 246, 134, 20, 11],
[2, 6, 120, 242, 65, 180, 227, 134, 14, 10],
[1, 4, 44, 120, 121, 11, 605, 79, 3, 12],
[14, 2, 55, 104, 65, 50, 141, 519, 11, 39]])