

```
In [1]: import torch
from torch import nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets
import torchvision.transforms as transforms
from torchvision.transforms import ToTensor
import time
from sklearn.model_selection import KFold
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import torch.nn.functional as F
import gc
from ptflops import get_model_complexity_info
from torchmetrics.classification import MulticlassConfusionMatrix
```

Problem 1

```
In [2]: data_unversioned/ecgr4106/'
s.CIFAR10(data_path, train=True, download=True, transform=transforms.Compose([transforms.ToTensor(), transforms.Resize(size=
assets.CIFAR10(data_path, train=False, download=True, transform=transforms.Compose([transforms.ToTensor(), transforms.Resize(
Files already downloaded and verified
Files already downloaded and verified
```

```
In [3]: def try_gpu(i=0):
    if torch.cuda.device_count() >= i+1:
        return torch.device(f'cuda:{i}')
    return torch.device('cpu')
```

```

In [4]: ▶ def training_loop(n_epochs, optimizer, model, loss_fn, train_loader, val_loader, update_freq):
    train_loss_hist = []
    train_acc_hist = []
    val_acc_hist = []
    main_tic = time.perf_counter()

    for epoch in range(1, n_epochs + 1):
        tic = time.perf_counter()
        loss_train = 0.0
        correct_train = 0
        correct_val = 0
        model_argmax = []
        labels_argmax = []

        for imgs, lbls in train_loader:
            images = imgs.to(device=try_gpu())
            labels = lbls.to(device=try_gpu())
            outputs = model(images)
            del images
            loss = loss_fn(outputs, labels)
            del labels
            del outputs
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            loss_train += loss.item()
            gc.collect()
            torch.cuda.empty_cache()

        toc = time.perf_counter()

        with torch.no_grad():
            total = 0
            for imgs, lbls in train_loader:
                images = imgs.to(device=try_gpu())
                labels = lbls.to(device=try_gpu())
                outputs = model(images)
                del images
                _, predicted = torch.max(outputs, dim=1)
                del outputs
                total += labels.shape[0]
                correct_train += int((predicted == labels).sum())
                del labels
                del predicted
            train_acc = round(correct_train/total, 3)
            total = 0
            for imgs, lbls in val_loader:
                images = imgs.to(device=try_gpu())
                labels = lbls.to(device=try_gpu())
                outputs = model(images)
                del images
                _, predicted = torch.max(outputs, dim=1)
                del outputs
                if epoch == 1 or epoch == n_epochs or epoch % update_freq == 0:
                    model_argmax = model_argmax + predicted.tolist()
                    labels_argmax = labels_argmax + labels.tolist()
                total += labels.shape[0]
                correct_val += int((predicted == labels).sum())
                del labels
                del predicted
            val_acc = round(correct_val/total, 3)

        train_loss_hist.append(round(loss_train / len(train_loader), 5))
        train_acc_hist.append(train_acc)
        val_acc_hist.append(val_acc)
        label_set = set(labels_argmax)

        if epoch == 1 or epoch == n_epochs or epoch % update_freq == 0:
            print(f"Epoch {epoch}: \n\tDuration = {round(toc - tic, 3)} seconds\n\tTraining Loss: {train_loss_hist[-1]}\n\tT
            metric = MulticlassConfusionMatrix(num_classes=len(label_set))
            print(metric(torch.ByteTensor(model_argmax), torch.ByteTensor(labels_argmax)))

    main_toc = time.perf_counter()
    print(f"\nTotal Training Time = {round(main_toc - main_tic, 3)} seconds\nAverage Training Time per Epoch (including val
    return train_loss_hist, train_acc_hist, val_acc_hist

```

```
In [5]: ▶ def plot_model(title, loss_hist, train_hist, test_hist, leg_loc):
fig, ax1 = plt.subplots()
x = range(1, len(loss_hist)+1)
ax1.plot(x, loss_hist, color='k')
ax1.set_xlabel('Epoch')
ax1.set_ylabel('Error')
ax1.tick_params(axis='y')

ax2 = ax1.twinx()
ax2.set_ylabel('Accuracy')
ax2.plot(x, train_hist)
ax2.plot(x, test_hist)
ax2.set_ylim([0, 1])
ax1.tick_params(axis='y')

fig.legend(["Training Loss", "Training Accuracy", "Testing Accuracy"], loc=leg_loc, bbox_to_anchor=(1, 1), bbox_transf
plt.title(title)
```

```
In [6]: ▶ def vgg_block(num_convs, out_channels):
layers = []
for _ in range(num_convs):
    layers.append(nn.LazyConv2d(out_channels, kernel_size=3, padding=1))
    layers.append(nn.LazyBatchNorm2d())
    layers.append(nn.ReLU())
layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
return nn.Sequential(*layers)
```

```
In [82]: ▶ def build_vgg(arch, num_classes=10):
conv_blks = []
for (num_convs, out_channels) in arch:
    conv_blks.append(vgg_block(num_convs, out_channels))
model = nn.Sequential(
    *conv_blks, nn.Flatten(),
    nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
    nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
    nn.LazyLinear(num_classes))
return model
```

```
In [ ]: ▶ train_loader_1 = DataLoader(cifar10, batch_size=64, shuffle=True)
val_loader_1 = DataLoader(cifar10_val, batch_size=64, shuffle=False)
```

```
In [83]: class tinyVGG(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(tinyVGG, self).__init__()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
        self.fc1 = nn.Linear(128)
        self.fc2 = nn.Linear(64)
        self.fc3 = nn.Linear(num_classes)
        self.fc_drop = nn.Dropout(p=0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.conv_blks(x)
        out = self.fc_drop(self.relu(self.fc1(out)))
        out = self.fc_drop(self.relu(self.fc2(out)))
        out = self.fc3(out)
        return out

model_0 = tinyVGG(arch=((1, 64), (1, 128))).to(device=try_gpu())
optimizer_0 = optim.SGD(model_0.parameters(), lr=0.1)
model_0.eval()
```

```
Out[83]: tinyVGG(
  (vgg): Sequential(
    (0): Sequential(
      (0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Flatten(start_dim=1, end_dim=-1)
    (3): LazyLinear(in_features=0, out_features=4096, bias=True)
    (4): ReLU()
    (5): Dropout(p=0.5, inplace=False)
    (6): LazyLinear(in_features=0, out_features=4096, bias=True)
    (7): ReLU()
    (8): Dropout(p=0.5, inplace=False)
    (9): LazyLinear(in_features=0, out_features=10, bias=True)
  )
)
```

```
In [ ]: with torch.cuda.device(0):
    macs, params = get_model_complexity_info(model_1, (1, 3, 64, 64), as_strings=True,
                                              print_per_layer_stat=False, verbose=False)
```

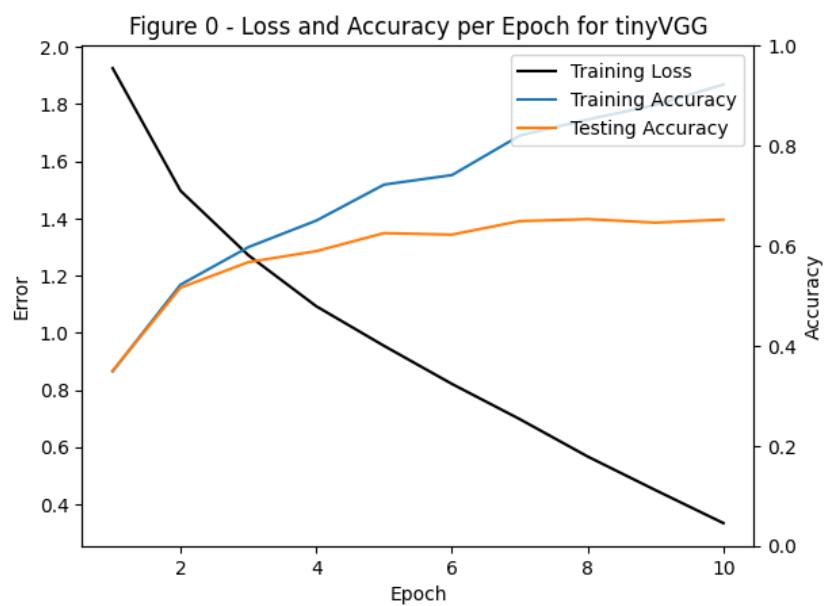
```
In [9]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_0, t_acc_hist_0, v_acc_hist_0 = training_loop(10,
    optimizer_0,
    model_0,
    nn.CrossEntropyLoss(),
    train_loader_1,
    val_loader_1,
    1)
```

```
[ 10, 17, 40, 49, 71, 37, 720, 15, 13, 28],
[ 20, 23, 38, 21, 59, 50, 8, 717, 5, 59],
[ 92, 119, 8, 5, 6, 7, 5, 5, 683, 70],
[ 12, 186, 6, 5, 3, 6, 6, 11, 17, 748]]

Epoch 9:
Duration = 163.986 seconds
Training Loss: 0.44921
Training Accuracy: 0.881
Validation Accuracy: 0.646
tensor([[766, 15, 67, 8, 9, 23, 5, 16, 24, 67],
        [ 29, 675, 28, 11, 4, 13, 10, 5, 18, 207],
        [ 78, 6, 647, 36, 61, 86, 39, 27, 7, 13],
        [ 22, 8, 149, 334, 33, 331, 66, 30, 15, 12],
        [ 28, 2, 186, 44, 513, 75, 75, 70, 4, 3],
        [ 12, 3, 110, 74, 38, 696, 25, 33, 4, 5],
        [ 8, 2, 91, 47, 43, 82, 703, 8, 4, 12],
        [ 15, 4, 89, 31, 49, 99, 11, 681, 4, 17],
        [168, 32, 31, 14, 7, 17, 7, 6, 665, 53],
        [ 37, 64, 22, 15, 6, 34, 5, 16, 20, 781]])

Epoch 10:
```

```
In [19]: title_0 = "Figure 0 - Loss and Accuracy per Epoch for tinyVGG"
plot_model(title_0, t_loss_hist_0, t_acc_hist_0, v_acc_hist_0, 'upper right')
```



```
In [84]: class VGG_11(nn.Module):
def __init__(self, arch, num_classes=10):
    super(VGG_11, self).__init__()
    self.vgg = build_vgg(arch)

def forward(self, x):
    return self.vgg(x)

model_1 = VGG_11(arch=((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))).to(device=try_gpu())
optimizer_1 = optim.SGD(model_1.parameters(), lr=0.1)
model_1.eval()
```

```
Out[84]: VGG_11(
  (vgg): Sequential(
    (0): Sequential(
      (0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Sequential(
      (0): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (3): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (4): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (5): Flatten(start_dim=1, end_dim=-1)
    (6): LazyLinear(in_features=0, out_features=4096, bias=True)
    (7): ReLU()
    (8): Dropout(p=0.5, inplace=False)
    (9): LazyLinear(in_features=0, out_features=4096, bias=True)
    (10): ReLU()
    (11): Dropout(p=0.5, inplace=False)
    (12): LazyLinear(in_features=0, out_features=10, bias=True)
  )
)
```

```
In [85]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_1, t_acc_hist_1, v_acc_hist_1 = training_loop(10,
                                                         optimizer_1,
                                                         model_1,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         2)
```

Epoch 1:

Duration = 306.019 seconds

Training Loss: 2.30314

Training Accuracy: 0.1

Validation Accuracy: 0.1

```
tensor([[ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0],
        [ 0,  0,  0, 1000,  0,  0,  0,  0,  0,  0]])
```

KeyboardInterrupt Traceback (most recent call last)

~\AppData\Local\Temp\ipykernel_11660\3450204939.py in <module>

1 torch.cuda.empty_cache()

2 gc.collect()

```
----> 3 t_loss_hist_1, t_acc_hist_1, v_acc_hist_1 = training_loop(10,
4                                                         optimizer_1,
5                                                         model_1,
```

~\AppData\Local\Temp\ipykernel_11660\310723436.py in training_loop(n_epochs, optimizer, model, loss_fn, train_loader, val_loader, update_freq)

24 loss.backward()

25 optimizer.step()

```
----> 26 loss_train += loss.item()
```

27 gc.collect()

28 torch.cuda.empty_cache()

KeyboardInterrupt:

```
In [105]: torch.cuda.empty_cache()
gc.collect()
```

Out[105]: 0

```
In [ ]: title_1 = "Figure 1 - Loss and Accuracy per Epoch for VGG-11"
plot_model(title_1, t_loss_hist_1, t_acc_hist_1, v_acc_hist_1, 'upper right')
```

```
In [8]: class VGG_16(nn.Module):
def __init__(self, arch, num_classes=10):
    super(VGG_16, self).__init__()
    self.vgg = build_VGG(arch)

def forward(self, x):
    return self.vgg(x)

model_2 = VGG_16(arch=((2, 64), (2, 128), (3, 256), (3, 512), (3, 512))).to(device=try_gpu())
optimizer_2 = optim.SGD(model_2.parameters(), lr=0.1)
model_2.eval()
```

C:\Users\ccm51\anaconda3\lib\site-packages\torch\nn\modules\lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to the API or functionality can happen at any moment.
warnings.warn('Lazy modules are a new feature under heavy development ')

```
Out[8]: VGG_16(
  (conv_blks): Sequential(
    (0): Sequential(
      (0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Sequential(
      (0): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (3): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (4): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (5): Flatten(start_dim=1, end_dim=-1)
  )
  (fc1): LazyLinear(in_features=0, out_features=4096, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=4096, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (fc_drop): Dropout(p=0.5, inplace=False)
  (relu): ReLU()
)
```



```
In [ ]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_2, t_acc_hist_2, v_acc_hist_2 = training_loop(3,
                                                         optimizer_2,
                                                         model_2,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)
```

```
In [ ]: title_2 = "Figure 2 - Loss and Accuracy per Epoch for VGG-16"
plot_model(title_2, t_loss_hist_2, t_acc_hist_2, v_acc_hist_2, 'upper right')
```

```
In [54]: class VGG_19(nn.Module):
    def __init__(self, arch, num_classes=10):
        super(VGG_19, self).__init__()
        conv_blks = []
        for (num_convs, out_channels) in arch:
            conv_blks.append(vgg_block(num_convs, out_channels))
        self.conv_blks = nn.Sequential(*conv_blks, nn.Flatten())
        self.fc1 = nn.Linear(4096)
        self.fc2 = nn.Linear(4096)
        self.fc3 = nn.Linear(num_classes)
        self.fc_drop = nn.Dropout(p=0.5)
        self.relu = nn.ReLU()

    def forward(self, x):
        out = self.conv_blks(x)
        out = self.fc_drop(self.relu(self.fc1(out)))
        out = self.fc_drop(self.relu(self.fc2(out)))
        out = self.fc3(out)
        return out

model_3 = VGG_19(arch=((2, 64), (2, 128), (4, 256), (4, 512), (4, 512))).to(device=try_gpu())
optimizer_3 = optim.SGD(model_3.parameters(), lr=0.1)
model_3.eval()
```

```

Out[54]: VGG_19(
  (conv_blks): Sequential(
    (0): Sequential(
      (0): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (1): Sequential(
      (0): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (2): Sequential(
      (0): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (10): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (11): ReLU()
      (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (3): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (10): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (11): ReLU()
      (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (4): Sequential(
      (0): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU()
      (3): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (5): ReLU()
      (6): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (8): ReLU()
      (9): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (10): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (11): ReLU()
      (12): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (5): Flatten(start_dim=1, end_dim=-1)
  )
  (fc1): LazyLinear(in_features=0, out_features=4096, bias=True)
  (fc2): LazyLinear(in_features=0, out_features=4096, bias=True)
  (fc3): LazyLinear(in_features=0, out_features=10, bias=True)
  (fc_drop): Dropout(p=0.5, inplace=False)
  (relu): ReLU()
)

```

```
In [ ]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_3, t_acc_hist_3, v_acc_hist_3 = training_loop(10,
                                                         optimizer_3,
                                                         model_3,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         2)

In [ ]: title_3 = "Figure 3 - Loss and Accuracy per Epoch for VGG-19"
plot_model(title_3, t_loss_hist_3, t_acc_hist_3, v_acc_hist_3, 'upper right')
```

Problem 2

```
In [11]: class Inception(nn.Module):
def __init__(self, c1, c2, c3, c4, **kwargs):
    super(Inception, self).__init__(**kwargs)
    self.b1_1 = nn.LazyConv2d(c1, kernel_size=1)
    self.b2_1 = nn.LazyConv2d(c2[0], kernel_size=1)
    self.b2_2 = nn.LazyConv2d(c2[1], kernel_size=3, padding=1)
    self.b3_1 = nn.LazyConv2d(c3[0], kernel_size=1)
    self.b3_2 = nn.LazyConv2d(c3[1], kernel_size=5, padding=2)
    self.b4_1 = nn.MaxPool2d(kernel_size=3, stride=1, padding=1)
    self.b4_2 = nn.LazyConv2d(c4, kernel_size=1)

def forward(self, x):
    b1 = F.relu(self.b1_1(x))
    b2 = F.relu(self.b2_2(F.relu(self.b2_1(x))))
    b3 = F.relu(self.b3_2(F.relu(self.b3_1(x))))
    b4 = F.relu(self.b4_2(self.b4_1(x)))
    return torch.cat((b1, b2, b3, b4), dim=1)
```

```

In [12]: class GoogLeNet(nn.Module):
    def __init__(self, num_classes=10):
        super(GoogLeNet, self).__init__()
        self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
                                   nn.LazyConv2d(64, kernel_size=1),
                                   nn.ReLU(),
                                   nn.LazyConv2d(192, kernel_size=3, padding=1),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body1 = nn.Sequential(Inception(64, (96, 128), (16, 32), 32),
                                    Inception(128, (128, 192), (32, 96), 64),
                                    nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body2 = nn.Sequential(Inception(192, (96, 208), (16, 48), 64),
                                    Inception(160, (112, 124), (24, 64), 64),
                                    Inception(128, (128, 256), (24, 64), 64),
                                    Inception(112, (144, 288), (32, 64), 64),
                                    Inception(256, (160, 320), (32, 128), 128),
                                    nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body3 = nn.Sequential(Inception(256, (160, 320), (32, 128), 128),
                                    Inception(384, (192, 384), (48, 128), 128),
                                    nn.AdaptiveAvgPool2d((1,1)),
                                    nn.Flatten())
        self.fc = nn.LazyLinear(num_classes)

    def forward(self, x):
        out = self.stem(x)
        out = self.body1(out)
        out = self.body2(out)
        out = self.body3(out)
        out = self.fc(out)
        return out

model_4 = GoogLeNet().to(device=try_gpu())
optimizer_4 = optim.SGD(model_4.parameters(), lr=0.1)
model_4.eval()

```

```

Out[12]: GoogleNet(
  (stem): Sequential(
    (0): LazyConv2d(0, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (3): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    (4): ReLU()
    (5): LazyConv2d(0, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU()
    (7): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  )
  (body1): Sequential(
    (0): Inception(
      (b1_1): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 96, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 16, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
    )
    (1): Inception(
      (b1_1): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 96, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    )
    (2): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  )
  (body2): Sequential(
    (0): Inception(
      (b1_1): LazyConv2d(0, 192, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 96, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 208, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 16, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 48, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    )
    (1): Inception(
      (b1_1): LazyConv2d(0, 160, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 112, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 124, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 24, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    )
    (2): Inception(
      (b1_1): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 24, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): Inception(
      (b1_1): LazyConv2d(0, 112, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 144, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 288, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 64, kernel_size=(1, 1), stride=(1, 1))
    )
    (4): Inception(
      (b1_1): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 160, kernel_size=(1, 1), stride=(1, 1))
      (b2_2): LazyConv2d(0, 320, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (b3_1): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
      (b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
      (b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
    )
    (5): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  )
  (body3): Sequential(
    (0): Inception(
      (b1_1): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(1, 1))
      (b2_1): LazyConv2d(0, 160, kernel_size=(1, 1), stride=(1, 1))

```

```

(b2_2): LazyConv2d(0, 320, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(b3_1): LazyConv2d(0, 32, kernel_size=(1, 1), stride=(1, 1))
(b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
(b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
(b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
)
(1): Inception(
  (b1_1): LazyConv2d(0, 384, kernel_size=(1, 1), stride=(1, 1))
  (b2_1): LazyConv2d(0, 192, kernel_size=(1, 1), stride=(1, 1))
  (b2_2): LazyConv2d(0, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (b3_1): LazyConv2d(0, 48, kernel_size=(1, 1), stride=(1, 1))
  (b3_2): LazyConv2d(0, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (b4_1): MaxPool2d(kernel_size=3, stride=1, padding=1, dilation=1, ceil_mode=False)
  (b4_2): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
)
(2): AdaptiveAvgPool2d(output_size=(1, 1))
(3): Flatten(start_dim=1, end_dim=-1)
)
(fc): LazyLinear(in_features=0, out_features=10, bias=True)
)

```

```

In [13]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_4, t_acc_hist_4, v_acc_hist_4 = training_loop(3,
                                                         optimizer_4,
                                                         model_4,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)

```

Epoch 1:

Duration = 243.455 seconds
 Training Loss: 2.30299
 Training Accuracy: 0.1
 Validation Accuracy: 0.1

```

tensor([[ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0],
        [ 0,  0,  0,  0,  0, 1000,  0,  0,  0,  0]])

```

Epoch 2:

Duration = 264.773 seconds
 Training Loss: 2.30299
 Training Accuracy: 0.1
 Validation Accuracy: 0.1

```

tensor([[ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0],
        [ 0,  0, 1000,  0,  0,  0,  0,  0,  0,  0]])

```

Epoch 3:

Duration = 266.93 seconds
 Training Loss: 2.30306
 Training Accuracy: 0.1
 Validation Accuracy: 0.1

```

tensor([[ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0],
        [ 0,  0,  0,  0, 1000,  0,  0,  0,  0,  0]])

```

Total Training Time = 1035.319 seconds

Average Training Time per Epoch (including validation) = 345.106 seconds

```

In [ ]: title_4 = "Figure 4 - Loss and Accuracy per Epoch for GoogLeNet"
plot_model(title_4, t_loss_hist_4, t_acc_hist_4, v_acc_hist_4, 'upper right')

In [ ]: class AltGoogLeNet(nn.Module):
    def __init__(self, num_classes=10):
        super(GoogLeNet, self).__init__()
        self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
                                   nn.LazyConv2d(64, kernel_size=1),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.LazyConv2d(192, kernel_size=3, padding=1),
                                   nn.LazyBatchNorm2d(),
                                   nn.ReLU(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body1 = nn.Sequential(Inception(64, (96, 128), (16, 32), 32),
                                   Inception(128, (128, 192), (32, 96), 64),
                                   nn.LazyBatchNorm2d(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body2 = nn.Sequential(Inception(192, (96, 208), (16, 48), 64),
                                   Inception(160, (112, 124), (24, 64), 64),
                                   Inception(128, (128, 256), (24, 64), 64),
                                   Inception(112, (144, 288), (32, 64), 64),
                                   Inception(256, (160, 320), (32, 128), 128),
                                   nn.LazyBatchNorm2d(),
                                   nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
        self.body3 = nn.Sequential(Inception(256, (160, 320), (32, 128), 128),
                                   Inception(384, (192, 384), (48, 128), 128),
                                   nn.LazyBatchNorm2d(),
                                   nn.AdaptiveAvgPool2d((1,1)),
                                   nn.Flatten())
        self.fc = nn.LazyLinear(num_classes)

    def forward(self, x):
        out = self.stem(x)
        out = self.body1(out)
        out = self.body2(out)
        out = self.body3(out)
        out = self.fc(out)
        return out

model_5 = GoogLeNet().to(device=try_gpu())
optimizer_5 = optim.SGD(model_5.parameters(), lr=0.1)
model_5.eval()

```

```

In [ ]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_5, t_acc_hist_5, v_acc_hist_5 = training_loop(5,
                                                         optimizer_5,
                                                         model_5,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         3)

```

```

In [ ]: title_5 = "Figure 5 - Loss and Accuracy per Epoch for GoogLeNet with Batch Norm"
plot_model(title_5, t_loss_hist_5, t_acc_hist_5, v_acc_hist_5, 'upper right')

```

Problem 3


```
In [20]: class Residual(nn.Module):
    def __init__(self, num_channels, use_1x1conv=False, strides=1):
        super().__init__()
        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1, stride=strides)
        self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1)
        if use_1x1conv:
            self.conv3 = nn.LazyConv2d(num_channels, kernel_size=1, stride=strides)
        else:
            self.conv3 = None
        self.bn = nn.LazyBatchNorm2d()

    def forward(self, x):
        y = F.relu(self.bn(self.conv1(x)))
        y = self.bn(self.conv2(y))
        if self.conv3:
            x = self.bn(self.conv3(x))
        y += x
        return F.relu(y)
```

```
In [48]: def block(num_residuals, num_channels, first_block=False):
    blk = []
    for i in range(num_residuals):
        if i == 0 and not first_block:
            blk.append(Residual(num_channels, use_1x1conv=True, strides=2))
        else:
            blk.append(Residual(num_channels))
    return nn.Sequential(*blk)
```

```
model_6 = ResNet(arch=((2, 64), (2, 128), (2, 256), (2, 512))).to(device=try_gpu())
optimizer_6 = optim.SGD(model_6.parameters(), lr=0.01)
model_6.eval()
```

```

Out[51]: ResNet(
  (stem): Sequential(
    (0): LazyConv2d(0, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
    (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  )
  (blks): Sequential(
    (0): Sequential(
      (0): Residual(
        (conv1): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (1): Residual(
        (conv1): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Sequential(
      (0): Residual(
        (conv1): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (conv2): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv3): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(2, 2))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (1): Residual(
        (conv1): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (2): Sequential(
      (0): Residual(
        (conv1): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (conv2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv3): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(2, 2))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (1): Residual(
        (conv1): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (3): Sequential(
      (0): Residual(
        (conv1): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (conv2): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv3): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(2, 2))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (1): Residual(
        (conv1): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (head): Sequential(
    (0): AdaptiveAvgPool2d(output_size=(1, 1))
    (1): Flatten(start_dim=1, end_dim=-1)
    (2): LazyLinear(in_features=0, out_features=10, bias=True)
  )
)

```

```
In [53]: ▶ torch.cuda.empty_cache()
gc.collect()
t_loss_hist_6, t_acc_hist_6, v_acc_hist_6 = training_loop(10,
                                                         optimizer_6,
                                                         model_6,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         2)
```

```
Epoch 1:
  Duration = 215.568 seconds
  Training Loss: 1.87298
  Training Accuracy: 0.324
  Validation Accuracy: 0.334
  tensor([[480, 101, 26, 85, 7, 54, 26, 71, 94, 56],
         [ 97, 462, 10, 47, 3, 92, 22, 56, 101, 110],
         [ 75, 35, 78, 119, 193, 113, 245, 121, 10, 11],
         [ 17, 28, 23, 233, 53, 181, 307, 135, 12, 11],
         [ 42, 15, 23, 74, 345, 71, 299, 112, 11, 8],
         [ 15, 22, 25, 169, 64, 274, 291, 123, 8, 9],
         [ 4, 10, 23, 102, 232, 63, 468, 92, 2, 4],
         [ 26, 34, 28, 86, 82, 80, 175, 456, 5, 28],
         [313, 128, 12, 74, 0, 63, 10, 47, 278, 75],
         [ 84, 161, 13, 68, 3, 68, 27, 196, 110, 270]])

Epoch 2:
  Duration = 237.357 seconds
  Training Loss: 1.80335
  Training Accuracy: 0.349
  Validation Accuracy: 0.347
  tensor([[422, 87, 3, 25, 1, 40, 4, 24, 319, 75],
         [ 56, 466, 0, 10, 0, 35, 3, 10, 251, 169],
         [107, 64, 27, 88, 48, 282, 163, 115, 66, 40],
         [ 49, 53, 8, 215, 5, 378, 82, 68, 64, 78],
         [ 52, 28, 4, 82, 102, 244, 251, 134, 58, 45],
         [ 31, 57, 6, 142, 11, 463, 65, 87, 79, 59],
         [ 12, 58, 2, 140, 26, 276, 333, 86, 24, 43],
         [ 46, 78, 5, 93, 18, 172, 37, 351, 36, 164],
         [144, 86, 1, 18, 0, 30, 0, 9, 623, 89],
         [ 54, 150, 0, 25, 1, 31, 2, 25, 246, 466]])

Epoch 4:
  Duration = 248.569 seconds
  Training Loss: 1.6383
  Training Accuracy: 0.42
  Validation Accuracy: 0.421
  tensor([[413, 94, 17, 35, 3, 5, 32, 22, 323, 56],
         [ 16, 653, 2, 31, 0, 5, 29, 10, 119, 135],
         [ 82, 48, 137, 178, 116, 30, 259, 79, 49, 22],
         [ 31, 46, 37, 430, 18, 45, 259, 39, 51, 44],
         [ 53, 29, 49, 151, 191, 20, 336, 109, 51, 11],
         [ 20, 40, 40, 390, 28, 142, 197, 66, 49, 28],
         [ 7, 33, 18, 152, 25, 4, 691, 21, 22, 27],
         [ 39, 46, 24, 185, 39, 27, 104, 431, 32, 73],
         [ 81, 121, 7, 36, 1, 5, 11, 11, 671, 56],
         [ 27, 259, 6, 44, 0, 1, 36, 21, 153, 453]])

Epoch 6:
  Duration = 237.792 seconds
  Training Loss: 1.51366
  Training Accuracy: 0.486
  Validation Accuracy: 0.489
  tensor([[526, 55, 68, 21, 16, 19, 16, 27, 175, 77],
         [ 36, 647, 11, 21, 1, 20, 17, 18, 61, 168],
         [ 58, 30, 411, 60, 140, 99, 90, 72, 21, 19],
         [ 28, 27, 105, 304, 58, 226, 135, 66, 17, 34],
         [ 49, 13, 178, 55, 366, 75, 109, 121, 18, 16],
         [ 12, 15, 121, 146, 65, 436, 80, 84, 22, 19],
         [ 3, 15, 83, 119, 127, 59, 535, 35, 5, 19],
         [ 30, 17, 45, 70, 71, 114, 40, 555, 10, 48],
         [160, 85, 35, 40, 13, 7, 6, 18, 543, 93],
         [ 37, 195, 16, 32, 1, 22, 21, 41, 71, 564]])

Epoch 8:
  Duration = 271.673 seconds
  Training Loss: 1.4083
  Training Accuracy: 0.506
  Validation Accuracy: 0.5
  tensor([[544, 62, 74, 15, 13, 47, 38, 36, 75, 96],
         [ 15, 627, 16, 18, 1, 41, 23, 24, 15, 220],
         [ 46, 17, 340, 83, 88, 163, 146, 79, 13, 25],
         [ 9, 12, 46, 339, 24, 297, 156, 72, 6, 39],
         [ 27, 9, 101, 98, 234, 145, 227, 139, 4, 16],
         [ 7, 7, 55, 147, 19, 561, 92, 87, 4, 21],
         [ 3, 6, 37, 114, 32, 71, 690, 28, 1, 18],
         [ 12, 9, 23, 71, 38, 167, 51, 586, 5, 38],
         [185, 93, 34, 38, 10, 36, 20, 21, 416, 147],
         [ 33, 148, 11, 27, 2, 37, 32, 36, 12, 662]])

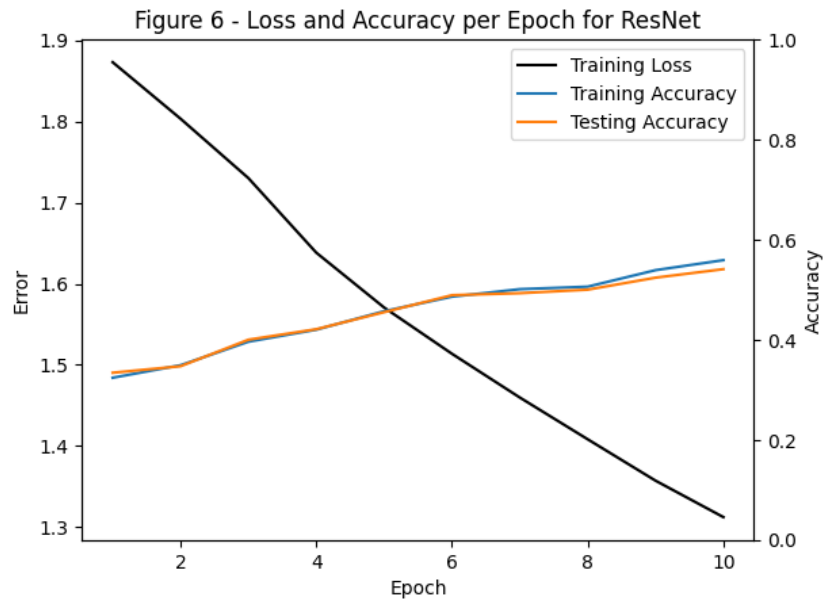
Epoch 10:
  Duration = 255.405 seconds
  Training Loss: 1.31234
  Training Accuracy: 0.559
  Validation Accuracy: 0.541
  tensor([[659, 40, 44, 11, 31, 15, 12, 19, 114, 55],
         [ 37, 741, 8, 6, 3, 13, 5, 5, 27, 155],
         [ 83, 28, 377, 22, 202, 126, 56, 59, 22, 25],
         [ 28, 23, 74, 172, 102, 336, 129, 71, 22, 43],
```

```
[ 51, 16, 95, 24, 515, 92, 73, 109, 13, 12],
[ 20, 14, 78, 50, 81, 561, 61, 96, 17, 22],
[ 18, 18, 56, 42, 160, 67, 578, 31, 8, 22],
[ 40, 13, 25, 25, 88, 141, 27, 585, 12, 44],
[190, 85, 24, 15, 10, 19, 6, 13, 589, 49],
[ 60, 187, 12, 5, 9, 17, 12, 28, 40, 630]]
```

Total Training Time = 3276.498 seconds

Average Training Time per Epoch (including validation) = 327.65 seconds

```
In [55]: title_6 = "Figure 6 - Loss and Accuracy per Epoch for ResNet"
plot_model(title_6, t_loss_hist_6, t_acc_hist_6, v_acc_hist_6, 'upper right')
```




```

In [102]: class ResNet26(nn.Module):
def __init__(self, num_classes=10):
    super(ResNet26, self).__init__()
    self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
    self.b1_1 = nn.Sequential(nn.LazyConv2d(64, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(64, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(256, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b1_2 = nn.Sequential(nn.LazyConv2d(256, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b2 = nn.Sequential(nn.LazyConv2d(64, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(64, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(256, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b3_1 = nn.Sequential(nn.LazyConv2d(128, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(128, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(512, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b3_2 = nn.Sequential(nn.LazyConv2d(512, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b4 = nn.Sequential(nn.LazyConv2d(128, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(128, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(512, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b5_1 = nn.Sequential(nn.LazyConv2d(256, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(256, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(1024, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b5_2 = nn.Sequential(nn.LazyConv2d(1024, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b6 = nn.Sequential(nn.LazyConv2d(256, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(256, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(1024, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b7_1 = nn.Sequential(nn.LazyConv2d(512, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(512, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(2048, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b7_2 = nn.Sequential(nn.LazyConv2d(2048, kernel_size=1),
                              nn.LazyBatchNorm2d())
    self.b8 = nn.Sequential(nn.LazyConv2d(512, kernel_size=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(512, kernel_size=3, padding=1),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.LazyConv2d(2048, kernel_size=1),
                              nn.LazyBatchNorm2d())

    self.relu = nn.ReLU()
    self.head = nn.Sequential(nn.AdaptiveAvgPool2d((1, 1)),
                              nn.Flatten(),
                              nn.LazyLinear(num_classes))

```



```

def forward(self, x):
    out = self.stem(x)
    out = self.relu(self.b1_1(out) + self.b1_2(out))
    out = self.relu(out + self.b2(out))
    out = self.relu(self.b3_1(out) + self.b3_2(out))
    out = self.relu(out + self.b4(out))
    out = self.relu(self.b5_1(out) + self.b5_2(out))
    out = self.relu(out + self.b6(out))
    out = self.relu(self.b7_1(out) + self.b7_2(out))
    out = self.relu(out + self.b8(out))
    out = self.head(out)
    return out

model_7 = ResNet26().to(device=try_gpu())
optimizer_7 = optim.SGD(model_7.parameters(), lr=0.01)
model_7.eval()

(1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(2): ReLU()
(3): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(5): ReLU()
(6): LazyConv2d(0, 256, kernel_size=(1, 1), stride=(1, 1))
(7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(b3_1): Sequential(
  (0): LazyConv2d(0, 128, kernel_size=(1, 1), stride=(1, 1))
  (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (2): ReLU()
  (3): LazyConv2d(0, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (5): ReLU()
  (6): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(1, 1))
  (7): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(b3_2): Sequential(
  (0): LazyConv2d(0, 512, kernel_size=(1, 1), stride=(1, 1))

```

```
In [106]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_7, t_acc_hist_7, v_acc_hist_7 = training_loop(3,
                                                         optimizer_7,
                                                         model_7,
                                                         nn.CrossEntropyLoss(),
                                                         train_loader_1,
                                                         val_loader_1,
                                                         1)
```

```
-----
OutOfMemoryError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11660\179788224.py in <module>
      1 torch.cuda.empty_cache()
      2 gc.collect()
----> 3 t_loss_hist_7, t_acc_hist_7, v_acc_hist_7 = training_loop(3,
      4                                                         optimizer_7,
      5                                                         model_7,

~\AppData\Local\Temp\ipykernel_11660\310723436.py in training_loop(n_epochs, optimizer, model, loss_fn, train_loader, val_loader, update_freq)
     16         images = imgs.to(device=try_gpu())
     17         labels = lbls.to(device=try_gpu())
--> 18         outputs = model(images)
     19         del images
     20         loss = loss_fn(outputs, labels)

~\anaconda3\lib\site-packages\torch\nn\modules\module.py in _call_impl(self, *input, **kwargs)
    1188         if not (self._backward_hooks or self._forward_hooks or self._forward_pre_hooks or _global_backward_hooks
    1189                 or _global_forward_hooks or _global_forward_pre_hooks):
-> 1190             return forward_call(*input, **kwargs)
    1191         # Do not call functions when jit is used
    1192         full_backward_hooks, non_full_backward_hooks = [], []

~\AppData\Local\Temp\ipykernel_11660\2953130020.py in forward(self, x)
     86         out = self.stem(x)
     87         out = self.relu(self.b1_1(out) + self.b1_2(out))
--> 88         out = self.relu(out + self.b2(out))
     89         out = self.relu(self.b3_1(out) + self.b3_2(out))
     90         out = self.relu(out + self.b4(out))

~\anaconda3\lib\site-packages\torch\nn\modules\module.py in _call_impl(self, *input, **kwargs)
    1188         if not (self._backward_hooks or self._forward_hooks or self._forward_pre_hooks or _global_backward_hooks
    1189                 or _global_forward_hooks or _global_forward_pre_hooks):
-> 1190             return forward_call(*input, **kwargs)
    1191         # Do not call functions when jit is used
    1192         full_backward_hooks, non_full_backward_hooks = [], []

~\anaconda3\lib\site-packages\torch\nn\modules\container.py in forward(self, input)
     202         def forward(self, input):
     203             for module in self:
--> 204                 input = module(input)
     205             return input
     206

~\anaconda3\lib\site-packages\torch\nn\modules\module.py in _call_impl(self, *input, **kwargs)
    1188         if not (self._backward_hooks or self._forward_hooks or self._forward_pre_hooks or _global_backward_hooks
    1189                 or _global_forward_hooks or _global_forward_pre_hooks):
-> 1190             return forward_call(*input, **kwargs)
    1191         # Do not call functions when jit is used
    1192         full_backward_hooks, non_full_backward_hooks = [], []

~\anaconda3\lib\site-packages\torch\nn\modules\batchnorm.py in forward(self, input)
     169         used for normalization (i.e. in eval mode when buffers are not None).
     170         """
--> 171         return F.batch_norm(
     172             input,
     173             # If buffers are not to be tracked, ensure that they won't be updated

~\anaconda3\lib\site-packages\torch\nn\functional.py in batch_norm(input, running_mean, running_var, weight, bias, training, momentum, eps)
    2448         _verify_batch_size(input.size())
    2449
-> 2450     return torch.batch_norm(
    2451         input, weight, bias, running_mean, running_var, training, momentum, eps, torch.backends.cudnn.enabled
    2452     )

OutOfMemoryError: CUDA out of memory. Tried to allocate 20.00 MiB (GPU 0; 2.00 GiB total capacity; 1.12 GiB already allocated; 0 bytes free; 1.15 GiB reserved in total by PyTorch) If reserved memory is >> allocated memory try setting max_split_size_mb to avoid fragmentation.  See documentation for Memory Management and PYTORCH_CUDA_ALLOC_CONF
```

```
In [ ]: title_7 = "Figure 7 - Loss and Accuracy per Epoch for ResNet26"
plot_model(title_7, t_loss_hist_7, t_acc_hist_7, v_acc_hist_7, 'upper right')
```

```
In [86]: class ResNet34(nn.Module):
def __init__(self, arch, num_classes=10):
    super(ResNet34, self).__init__()
    self.stem = nn.Sequential(nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                              nn.LazyBatchNorm2d(),
                              nn.ReLU(),
                              nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

    blks = []
    for i, b in enumerate(arch):
        blks.append(block(*b, first_block=(i==0)))
    self.blks = nn.Sequential(*blks)
    self.head = nn.Sequential(nn.AdaptiveAvgPool2d((1, 1)),
                              nn.Flatten(),
                              nn.LazyLinear(num_classes))

def forward(self, x):
    out = self.stem(x)
    out = self.blks(out)
    out = self.head(out)
    return out

model_8 = ResNet34(arch=((3, 64), (4, 128), (6, 256), (3, 512))).to(device=try_gpu())
optimizer_8 = optim.SGD(model_8.parameters(), lr=0.01)
model_8.eval()
```

```
Out[86]: ResNet34(
  (stem): Sequential(
    (0): LazyConv2d(0, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
    (1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  )
  (blks): Sequential(
    (0): Sequential(
      (0): Residual(
        (conv1): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    (1): Residual(
      (conv1): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (conv2): LazyConv2d(0, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (bn1): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (bn2): LazyBatchNorm2d(0, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
```

```
In [ ]: torch.cuda.empty_cache()
gc.collect()
t_loss_hist_7, t_acc_hist_7, v_acc_hist_7 = training_loop(3,
    optimizer_7,
    model_7,
    nn.CrossEntropyLoss(),
    train_loader_1,
    val_loader_1,
    1)
```

```
In [ ]: title_8 = "Figure 8 - Loss and Accuracy per Epoch for ResNet34"
plot_model(title_8, t_loss_hist_8, t_acc_hist_8, v_acc_hist_8, 'upper right')
```